
HACKING C++

https://hackingcpp.com/cpp/beginners_guide.html

Diagnostics

Basics: Warnings & Tests

Terms And Techniques

- Warnings compiler messages hinting at (see below)
 potentially problematic runtime
 behavior / subtle pitfalls
 - Assertions statements for comparing and (see below)
 reporting expected and actual
 values of expressions
 - Testing compare actual and expected (see below)
 behavior of parts or entire prog
 - Code Coverage how much code is actually gcov, ...
 executed and/or tested
 - Static Analysis finds potential runtime problems ASAN, UBSAN. ...
 like undefined behavior by
 analyzing the source code
 - Dynamic Analysis finds potential problems like valgrind, ...
 memory leaks by running the
 actual program
 - Debugging step through code at runtime and (next up)
 inspect in-memory values
 - Profiling find out how much each function/
 loop/code block contributes to
 the total running time, memory
 consumption, ...
 - Micro Bench-
 marking small tests that measure the
 runtime of single a function or
 a block of statements/calls rather
 than the whole program
- Remember: Use Dedicated Types!
- to restrict input parameter values
 - to ensure validity of intermediate results
 - to guarantee validity of return values

Compiler Warnings

- Compiler Error = program not compilable
 - Compiler Warning = program compilable, but there is a problematic
 piece of code that might lead to runtime bugs
- gcc/clang Options
- Most important
 - Wall, -Wextra, -Wpedantic, -Wshadow, -Werror
 - fsanitize=undefined,address
 - Recommended Set (Production Level)
 - Wall -Wextra -Wpedantic -Wshadow -Wconversion -Werror
 - fsanitize=undefined,address -Wfloat-equal -Wformat-nonliteral
 - Wformat-security -Wformat-y2k -Wformat=2 -Wimport -Winvalid-pch

```
-Wlogical-op -Wmissing-declarations -Wmissing-field-initializers
-Wmissing-format-attribute -Wmissing-include-dirs -Wmissing-noreturn
-Wnested-externs -Wpacked -Wpointer-arith -Wredundant-decls
-Wstack-protector -Wstrict-null-sentinel -Wswitch-enum -Wundef
-Wwrite-strings
```

- High Performance / Low Memory / Security

Might be VERY noisy!

```
-Wdisabled-optimization -Wpadded -Wsign-conversion -Wsign-promo
-Wstrict-aliasing=2 -Wstrict-overflow=5 -Wunused -Wunused-parameter
```

Assertions

- Runtime Assertions

```
#include <cassert>
assert(bool_expression);
```

- Aborts the program if expresion yields false

- Use cases:

- check expected values/conditions at runtime
- verify preconditions (input values)
- verify invariants (e.g., intermediate states/results)
- verify postconditions (output/return values)

- Important:

- Runtime assertions should be deactivated in release builds to avoid any performance impact

- Example

```
#include <cassert>

double sqrt (double x) {
    assert ( x >= );
    ...
}
...
double r = sqrt(-2.3);
```

- Commas must be protected by parentheses

- assert

- is a preprocessor macro
- commas are interpreted as macro argument separator:

```
assert( min(1,2) == 1 );           // ERROR
assert((min(1,2) == 1 ));         // OK
```

- Messages

- Can be added with a custom macro (there is no standard way)

```
#define assertmsg(expr, msg) assert (((void)msg, expr))

assertmsg(1+2==2, "1 plus 1 must be 2");
```

- (De-)Activation - g++/clang

- Assertions are deactivated by defining preprocessor macro:

```
NDEBUG
```

- Example

- With compiler switch:

```
$ g++ -DNBEBUG ...
```

- Static Assertions C++11

```
static_assert(bool_constexpr, "message");  
static_assert(bool_constexpr); C++17
```

- Aborts compilation if a compile-time constant expression yields false

```
...  
using index_t = int;  
index_t constexpr DIMS = 1; // oops  
  
void foo () {  
    static_assert(DIMS > 1, "DIMS must be at least 2");  
    ...  
}  
  
index_t bar (...) {  
    static_assert(  
        std::numeric_limits<index_t>::is_integer &&  
        std::numeric_limits<index_t>::is_signed,  
        "index type must be a signed integer");  
    ...  
}
```

Testing

- Guidelines

- Use Assertions

- to check expectations/assumptions that are not already expressible/guaranteed by types:
 - expected values that are only available at runtime
 - preconditions (input values)
 - invariants (e.g., intermediate states/results)
 - postconditions (output/return values)
- Note:
 - Runtime assertions should be deactivated in release builds
 - to avoid any performance impact

- Write Tests

- as soon as the basic purpose and interface of a function or type is decided
 - faster development: less need for time-consuming logging and debugging sessions
 - easier performance tuning: one can continuously check if still correct
 - documentation: expectations/assumptions are written down in code

- Use A Testing Framework

- More convenient and less error-prone:
 - predefined checks, setup facilities, test runner, ...

- Beginners/Smaller Projects:

doctest

- very compact and self-documenting style
- easy setup: only include a single header
- very fast compilation

- Larger Projects:

Catch2

- same basic philosophy as doctest
- value generators for performing same test with different values
- micro benchmarking with timer, averaging, etc...
- slower to compile and slightly more complicated to set up than doctest

- doctest - Simple Test Case (from doctest tutorial)

```
#define DOCTEST_CONFIG_IMPLEMENT_WITH_MAIN
#include "doctest.h"
```

```
int factorial (int n) {
    if (n <= 1) return n;
    return factorial(n-1) * n;
}

TEST_CASE("testing factorial") {
    CHECK(factorial(0) == 1);
    CHECK(factorial(1) == 1);
    CHECK(factorial(2) == 2);
    CHECK(factorial(3) == 6);
    CHECK(factorial(10) == 3628800);
}
```

- The test fails, because the implementation of factorial doesn't handle the case of n = 0 properly

- doctest - Subcases (from doctest tutorial)

```
TEST_CASE("vectors can be sized and resized") {
    std::vector<int> v(5);

    REQUIRE(v.size() == 5);
    REQUIRE(v.capacity() >= 5);

    SUBCASE("push_back increases the size") {
        v.push_back(1);
        CHECK(v.size() == 6);
        CHECK(v.capacity() >= 6);
    }
    SUBCASE("reserve increases the capacity") {
        v.reserve(6);
        CHECK(v.size() == 5);
        CHECK(v.capacity() >= 6);
    }
}
```

- If CHECK fails:
 - test is marked as fails, but execution continues
- If REQUIRE fails:
 - execution stops

- Don't Use con/cout/cerr Directly!

- Bad: direct use of global I/O streams makes functions or types hard to test:

```
void bad_log (State const& s) { std::cout << ... }
```

- Good: In Functions: Pass Streams By Reference

```
struct State { std::string msg; ... };

void log (std::ostream& os, State const& s) { os << s.msg; }

TEST_CASE("State Log") {
```

```

    State s {"expected"};
    std::ostringstream oss;
    log(oss, s);
    CHECK(oss.str() == "expected");
}

```

- Good: Class Scope: Store Stream Pointers
 - But try to write types that are independent of streams or any other particular I/O method!

```

class Logger {
    std::ostream* os_;
    int count_;
public:
    explicit
    Logger (std::ostream* os): os_{os}, cout_{0} {}
    ...
    bool add (std::string_view msg) {
        if (!os_) return false;
        *os_ << cout_ << ": " << msg << '\n';
        ++cout_;
        return true;
    }
};

TEST_CASE("Logging") {
    std::ostringstream oss;
    Logger log {&oss};
    log.add("message");
    CHECK(oss.str() == "o: message\n");
}

```

Debugging With gdb

- gdb / Frontends

- | | |
|----------|---|
| - gdb | - Qt Creator can connect to gdb |
| - cgdb | - Visual Studio Code can connect to gdb |
| - gdbgui | - Vim -> :help Termdebug |

- Example

```
$ vi sum.cpp
```

```

#include <iostream>
#include <cstdlib>

int sum_up_to (int n) {
    if (n < 0) return 0;
    int sum = 0;
    for (int i = 0; i < n; ++i) {
        sum += i;
    }
    return sum;
}

int main (int argc, char* argv[]) {
    if (argc < 2) {
        std::cout << "Usage: " << argv[0] << " <n>\n";
        return 0;
    }
    int n = std::atoi(argv[1]);
    int sum = sum_up_to(n);
    std::cout << "result: " << sum << '\n';
}

```

- Compile For Debugging

```
$ g++ -g -o sum sum.cpp
```

- Start The Debugger

```
$ gdb sum
(gdb)
```

- Running Your Program

```
run <args>
```

```
(gdb) run 5
Starting program: /home/.../sum
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/\
libthread_db.so.1".
result: 10
[Inferior 1 (process 6122) exited normally]
```

- Setting Breakpoints

Command	Stop at
break 12	line number 12 in currently active source code file
break sum_up_to	first executable line of function sum_up_to within ALL source code files
break sum.cpp:7	line number 7 in file sum.cpp (if not in the same directory, use relative/full pathname)
break sum.cpp:main	first executable line of function main in file sum.cpp

- Stepping Through Your Program

```
$ cgdb sum
(gdb) break main
Breakpoint 1 at 0x123b: file debug-01-sum.cpp, line 14.
```

```
(gdb) run 5
Starting program: /home/.../sum 5
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/\
libthread_db.so.1".
```

```
Breakpoint 1, main (argc=2, argv=0x7fffffffdfb8) at sum.cpp:14
14      if (argc < 2) {
```

```
(gdb) next
18      int n = std::atoi(argv[1]);
```

```
(gdb) next
19      int sum = sum_up_to(n);
```

```
(gdb) step
sum_up_to (n=5) at debug-01-sum.cpp:5
5      if (n < 0) return 0;
```

... several times next

```
(gdb) print sum
```

```
$2 = 1
```

```
... several times next
```

```
(gdb) next
result: 10
21      }
```

- Conditional Breakpoints

```
break 20 if i == 2000
```

```
break 180 if i < 0
```

```
break test.cpp:34 if (x & y) == 1
```

```
break myfunc if i % (j+3) != 0
```

- Managing Breakpoints

info breakpoints	..show all breakpoints
delete	..delete all breakpoints
delete 1	..delete breakpoint number 1
clear	..delete breakpoint at the next instruction
disable 2	..disable breakpoint number 2
enable 2	..(re-)enable breakpoint number 2
save breakpoints <file>	..save breakpoints to file
source <file>	..load breakpoints from file

- Inspecting and Setting Variables

- monitor local variables

```
info locals
```

- print value of a variable / an expression

```
print x
print x + 2
```

- set value of variable

```
set x = 12
```

- Useful gdb Commands

<Tab>	..complete command or function name
run [<arg>...]	..run program (with cmd line arg(s))
break <loc>	..set breakpoint at beginning of function or at a specific line
step	..next instruction, step into function
next	..next instruction, step over function
jump <loc>	..jump to location (useful for exiting long/endless loops)
continue	..continue until next breakpoint or end of program
until <loc>	..continue until location (function/line)
finish	..finish (step out of) current function
print <expression>	..print value of expression, e.g. variable
info breakpoints	..list all breakpoints
info locals	..list local variables and their values
backtrace	..show call stack

- Address Sanitizer (ASan)

- g++, clang++
- detects memory corruption bugs
 - memory leaks
 - access to already freed memory
 - access to incorrect stack areas
 - ...
- instruments your code with additional instructions
 - roughly 70% runtime increase
 - roughly 3-fold increase in memory usage

- Example: Dereferencing Null Pointer

```
$ vi asan.cpp

#include <iostream>

int main () {
    int* p = nullptr;
    cout << "p = " << *p << '\n';
}
```

- Address Sanitizer in Action

```
$ g++ asan.cpp -o asan -fsanitize=address
$ ./asan
=====
==7498==ERROR: AddressSanitizer: SEGV on unknown address \
0x000000000000 (pc 0x55f255a73315 bp 0x7ffc50ab1d70 sp \
0x7ffc50ab1d60 T0)
==7498==The signal is caused by a READ memory access.
==7498==Hint: address points to the zero page.
#0 0x55f255a73315 in main (/home/.../asan+0x1315)
#1 0x7fd6af429d8f in __libc_start_call_main ../sysdeps/nptl/\
libc_start_call_main.h:58
#2 0x7fd6af429e3f in __libc_start_main_impl ../csu/\
libc-start.c:392
#3 0x55f255a731e4 in _start (/home/.../asan+0x11e4)

AddressSanitizer can not provide additional info.
SUMMARY: AddressSanitizer: SEGV (/home/.../asan+0x1315) in main
==7498==ABORTING
```

Undefined Behavior Sanitizer

- Undefined Behavior Sanitizer (UBSAN)

- clang++, g++
- detects many types of undefined behavior at runtime
 - dereferencing null pointers
 - reading from misaligned pointers
 - integer overflow
 - division by zero
 - ...
- instruments your code with additional instructions:
 - runtime increase in debug build ~25%

- Example: Signed Integer Overflow

```
$ vi ubsan.cpp

int main()
{
    int i = std::numeric_limits<int>::max();
```



```

        i += 1;          // overflow!
        cout << "i = " << i << '\n';
    }

```

```

$ clang++ ubsan.cpp -o ubsan
$ ./ubsan
i = -2147483648

```

- UBSAN in Action

```

$ clang ++ ubsan.cpp -o ubsan -fsanitize=undefined
$ ./ubsan
ubsan.cpp:6:4: runtime error: signed integer overflow: 2147483647 + 1 \
cannot be represented in type 'int'
SUMMARY: UndefinedBehaviorSanitizer: undefined-behavior ubsan.cpp:6:4 \
in
i = -2147483648

```

valgrind

- Valgrind

- detects common runtime errors
 - readint/writing feed memory or incorrect stack areas
 - using values before they have been initialized
 - incorrect freeing of memory, such as double freeing
 - misuse of functions for memory allocations
 - memory leaks - unintentional memory consumption often related to program logic flaws
 - leading to loss of memory pointers prior to deallocation

```
valgrind [options] ./program [program options]
```

- Options

```

--help
--tool=memcheck    ..check the memory of your program
--leak-check=full  ..see the details of leaked memory
--verbose

```

- Example

```

$ vi compare.cpp

#include <iostream>

bool f (int i) {
    return (i == 5)
}

int main () {
    int i;
    if ( f(i) ) {
        cout << "MATCH\n";
    }
}

$ g++ -g -o compare compare.cpp
$ valgrind ./compare
==8216== Memcheck, a memory error detector
==8216== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward \
et al.
==8216== Using Valgrind-3.18.1 and LibVEX; rerun with -h for \
copyright info
==8216== Command: ./compare

```

```

==8216==
==8216== Conditional jump or move depends on uninitialised value(s)
==8216==    at 0x1091B5: main (compare.cpp:10)
==8216==
==8216==
==8216== HEAP SUMMARY:
==8216==    in use at exit: 0 bytes in 0 blocks
==8216==    total heap usage: 1 allocs, 1 frees, 72,704 bytes
        allocated
==8216==
==8216== All heap blocks were freed -- no leaks are possible
==8216==
==8216== Use --track-origins=yes to see where uninitialised values
        come from
==8216== For lists of detected and suppressed errors, rerun with: -s
==8216== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from
        0)

```

- Example

```
$ vi readwrite.cpp
```

```

#include <iostream>
#include <vector>

int main () {
    constexpr int dim = 100;
    int i;
    std::vector<float> v(dim);

    for(i = 0; i < dim; ++i)
        v[i] = 0.0;

    float k = 2.0f;
    float sup = k;

    k = v[i];
    v[i] = sup;

    std::cout << "GOOD END\n";
}

```

```

$ g++ readwrite.cpp -g -o readwrite
$ valgrind --tool=memcheck ./readwrite
==8359== Memcheck, a memory error detector
==8359== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward
        et al.
==8359== Using Valgrind-3.18.1 and LibVEX; rerun with -h for
        copyright info
==8359== Command: ./readwrite
==8359==
==8359== Invalid read of size 4
==8359==    at 0x109315: main (readwrite.cpp:15)
==8359== Address 0x4ddfe10 is 0 bytes after a block of size 400
        alloc'd
==8359==    at 0x4849013: operator new(unsigned long) (in
        /usr/libexec/valgrind/vgpreload_memcheck-amd64-linux.so)
==8359== by 0x109BEB: __gnu_cxx::new_allocator<float>::allocate\
        (unsigned long, void const*) (new_allocator.h:127)
==8359== by 0x109AC7: std::allocator_traits<std::allocator\
        <float> >::allocate(std::allocator<float>&, unsigned
        long) (alloc_traits.h:464)
==8359== by 0x109A01: std::_Vector_base<float, std::allocator\
        <float> >::_M_allocate(unsigned long) (stl_vector.h:346)
==8359== by 0x10987C: std::_Vector_base<float, std::allocator\
        <float> >::_M_create_storage(unsigned long) \
        (stl_vector.h:361)
==8359== by 0x10968E: std::_Vector_base<float, std::allocator\

```

```
<float> >::_Vector_base(unsigned long, std::allocator\
<float> const&) (stl_vector.h:305)
==8359== by 0x1094E4: std::vector<float, std::allocator<float> \
>::vector(unsigned long, std::allocator<float> const&) \
(stl_vector.h:511)
==8359== by 0x1092AC: main (readwrite.cpp:7)
==8359==
==8359== Invalid write of size 4
==8359== at 0x109342: main (readwrite.cpp:16)
==8359== Address 0x4ddfe10 is 0 bytes after a block of size 400 \
alloc'd
==8359== at 0x4849013: operator new(unsigned long) (in /usr/\
libexec/valgrind/vgpreload_memcheck-amd64-linux.so)
==8359== by 0x109BEB: __gnu_cxx::new_allocator<float>::allocate\
(unsigned long, void const*) (new_allocator.h:127)
==8359== by 0x109AC7: std::allocator_traits<std::allocator\
<float> >::allocate(std::allocator<float>&, unsigned \
long) (alloc_traits.h:464)
==8359== by 0x109A01: std::_Vector_base<float, std::allocator\
<float> >::_M_allocate(unsigned long) (stl_vector.h:346)
==8359== by 0x10987C: std::_Vector_base<float, std::allocator\
<float> >::_M_create_storage(unsigned long) \
(stl_vector.h:361)
==8359== by 0x10968E: std::_Vector_base<float, std::allocator\
<float> >::_Vector_base(unsigned long, std::allocator\
<float> const&) (stl_vector.h:305)
==8359== by 0x1094E4: std::vector<float, std::allocator<float> \
>::vector(unsigned long, std::allocator<float> const&) \
(stl_vector.h:511)
==8359== by 0x1092AC: main (readwrite.cpp:7)
==8359==
GOOD END
==8359==
==8359== HEAP SUMMARY:
==8359== in use at exit: 0 bytes in 0 blocks
==8359== total heap usage: 3 allocs, 3 frees, 74,128 bytes \
allocated
==8359==
==8359== All heap blocks were freed -- no leaks are possible
==8359==
==8359== For lists of detected and suppressed errors, rerun with: -s
==8359== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 \
from 0)
```

- Valgrind Tools

--tool=memcheck	..leaks, invalid reads/writes detection
--tool=callgrind	..runtime profiling
--tool=cachgrind	..cache profiling
--tool=massif	..heap memory profiling
--tool=helgrind	..locking order violation detection
--tool=drd	..multithreading error detection