
HACKING C++
https://hackingcpp.com/cpp/beginners_guide.html

Standard Library - Part 1

Iterators

- Objects that "point to a location"
 - May point to a readable memory address/object
- Used to:
 - Iterate over container elements in a data-layout-agnostic way
 - Specify positions and ranges in containers
 - For insertion, deletion, ...
- Note:
 - The @name notation is used to denote an iterator object/parameter
 - @ i neither an allowed operator nor does it have any other meaning in C++

- Default Iterators

- Obtainable from standard containers, either:
 - (1) with member functions

```
container.begin()    -> @first_element
container.end()      -> @one_behind_last_element
```

- (2) with free-standing functions (C++11)

```
std::begin(container) -> @first_element
std::end(container)   -> @one_behind_last_element
```

- An iterator refers to a position in a container

```
vector<int> v {1,2,3,4,5,6,7};
auto i = begin(v);
auto e = end(v);
```

```

  1|2|3|4|5|6|7| |
    ^             ^
    i             e

```

- *i accesses the element at i's position

```
cout << *i;           // prints 1
cout << *(i+2);       // prints 3
cout << *e;           // UNDEFINED BEHAVIOR
```

- The 'end' iterator
 - Is only intended to be used as position indicator
 - Must not be used to access an element!

- Forward/backward

```
++i ..goes 1 toward end
--i ..goes 1 towards begin

i += 2 ..goes forward by 2
i -= 3 ..goes backwards by 3
```

- Example

```
i += 5;
if (i == end(v))
    cout << "at end";
```

- Reverse Iterators

- Obtainable from (many, but not all) standard containers, either:
 - (1) with member functions:

```
container.rbegin()    -> @last_element
container.rend()      -> @one_behind_last_element
```

- (2) with free-standing functions (C++11)

```
std::rbegin(container) -> @last_element
std::rend()            -> @one_behind_last_element
```

- A reverse iterator refers to a position in a container

```
vector<int> v {1,2,3,4,5,6,7};
auto i = rbegin(v);
auto e = rend(v);
```

```

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| ^ | | | | | | | ^ |
| e | | | | | | | i |
```

- *i accesses the element at i's position

```
cout << *i;           // prints 7
cout << *(i+2);       // prints 5
cout << *e;           // UNDEFINED BEHAVIOR
```

- The 'rend' iterator
 - Is only intended to be used as position indicator
 - Must not be used to access an element!

- Forward/backward

```
++i ..goes 1 toward begin
--i ..goes 1 towards end

i += 2 ..goes backwards by 2
i -= 3 ..goes forward by 3
```

- Example

```
i += 5;
if (i == rend(v))
    cout << "at end";
```

- Returning a corresponding normal (non-reverse) iterator

```
ri.base()
```

- Reverse position = normal position -1
- Normal position = reverse position +1

```
vector<int> v {1,2,3};
auto re = rbegin(v);
auto fw = re.base();
```

```
re    v
```

```

      |1|2|3| |
      ^
    fw

```

- Iterator-Based Loops

- Forward Direction

- + works for all standard sequence containers
- out-of-bounds access bugs possible
- verbose

```

std::vector<int> v {1,2,3,4,5,6};
for (auto i = begin(v); i != end(v); ++i) {
    cout << *i;
}

```

- Reverse Direction

- + works for all bidirectional containers
- out-of-bounds access bugs possible
- verbose

```

std::vector<int> v {1,2,3,4,5,6};
for (auto i = rbegin(v); i != rend(v); ++i) {
    cout << *i;
}

```

- Example: Swap Adjacent Pairs

```

void swap_adjacent_pairs (std::vector<int>& v) {
    if (v.size() < 2) return;
    for (auto i=begin(v), j=i+1, e=end(v); j < e; i+=2, j+=2) {
        std::swap(*i,*j);
    }
}

vector<int> v {1,2,3,4,5,6};          // 1 2 3 4 5 6
swap_adjacent_pairs(v);               // 2 1 4 3 6 5

```

- Iterator Ranges

= pair p,q of iterators

```

      begin(v)+3  begin(v)+6
      |p      q|
      |v      v|
v |1|2|3|4|5|6|7|8|9| |
  0 1 2 3 4 5 6 7 8
  ^           ^
begin(v)      end(v)

```

- IMPORTANT:

- The end-of-range iterator q points ONE BEHIND the last element in the range!

```
std::vector<int> v {0,1,2,3,4,5};
```

```

begin(v),      end(v)
      6 elements  |0|1|2|3|4|5| |
                  ^           ^

```

```

begin(v)+2, begin(v)+5
      3 elements  |0|1|2|3|4|5| |
                  ^           ^

```

```

begin(v)+2, begin(v)+3
      1 element   |0|1|2|3|4|5| |

```

```
begin(v)+2, begin(v)+2
empty range |0|1| 2 |3|4|5| |
            ^ ^
```

- Used for specifying ranges of elements to be:

- erased from a container

```
std::vector<int> v {1,2,3,4,5,6,7,8,9};
v.erase(begin(v)+3, begin(v)+6)
```

```
|1|2|3|4|5|6|7|8|9| -> |1|2|3|7|8|9|
```

- inserted into a container
- assigned to a container
- processed by a standard algorithm
- ...

- Element Count of Iterator Range

```
distance(@begin, @end) -> number of elements
           v           v
|0|1|2|3|4|5|6|7|8| -> 5
```

- Returns the size of an iterator range

```
distance(@range_begin, @element_in_range) -> index_of_element_in_range
```

- Example

```
#include <vector>
#include <iostream>
#include <algorithm>
#include <iterator> // std::distance

std::vector<int> {0,1,2,3,4,5,6,7,8};

// size of subrange
auto n = distance(begin(v)+2, begin(v)+7); // int n = 5

// size of entire container
auto m = distance(begin(v), end(v)); // int m = 9

std::vector<int> w {4,5,1,9,8};
// get index of smallest element in w:
auto argmin = distance(begin(w), min_element(begin(w),end(w)) );
// int argmin = 2
```

- TIP:

- Avoid using 'distance' with iterators into non-random access containers
 - e.g., std::list
 - The runtime will be proportional to the size of the input range!

Standard Sequence Containers

```
array<T,size>      |a|1|2|3|4|5|6|
                   fixed-size contiguous array;
```

```

vector<T>          v -> |1|2|3|4|5|6|

                    dynamic contiguous array;
                    amortized O(1) growth strategy;
                    C++'s "default" container;

deque<T>           d -> |1|2|<->|3|4|5|<->|6|

                    double-ended queue;
                    fast insert/erase at both ends;

list<T>            l ->  v-----+
                        |1|<->|2|<->|3|<->|4|<->|5|<->|end|
                        +-----^

                    doubly-linked list;
                    O(1) insert, erase & splicing;
                    in practice often slower than vector;

forward_list<T>    fl -> |1|->|2|->...->|6|->|end|

                    singly-linked list;
                    O(1) insert, erase & splicing;
                    needs less memory than 'list';
                    in practice often slower than vector;

```

Common Features

- Regularity: Copy, Assign, Compare

- All standard sequence containers are regular types:
 - deeply copyable
 - copying creates a new container object and copies all contained values
 - deeply assignable
 - all contained objects are copied from source to assignment target
 - deeply comparable
 - comparing two containers compares the contained objects
 - deeply ownership
 - destroying the container destroys all contained objects

- Example

```

std::vector<int> a {4,7,1,9};
std::vector<int> b {1,2,3};

bool equal1 = (a==b);          // false

b = a;                          // copy assignment -> b: 4 7 1 9
bool equal2 = (a==b);          // true

a[0] = 3;                       // a: 3 7 1 9; b: 4 8 1 9
bool equal3 = (a==b);          // false

// different ways of making exact copies,
// i.e., copy-constructing new containers:
std::vector<int> a2 {a};
std::vector<int> a3 (a);
std::vector<int> a4 = a;
auto a5 {a};
auto a6 (a);
auto a7 = a;

```

- Type Argument Deduction (C++17)

- As of C++17 the element type can be deduced from constructor calls

```
std::vector v {1, 2, 3, 4};           // std::vector<int>
std::vector v {1.f, 2.3f, 4.5f};     // std::vector<float>
std::vector v {1., 2.3, 4.6};        // std::vector<double>

struct p2 { int x; int y; };
std::vector v {p2{1,2}};             //std::vector<p2>
```

- Common Interface Parts

- Iterators for Forward Traversal

- can be obtained from all standard sequence containers either with:

- (1) member functions:

```
container.begin() -> @first_element
container.end()   -> @one_behind_last_element
```

- (2) free-standing functions (C++11):

```
std::begin(container) -> @first_element
std::end(container)   -> @one_behind_last_element
```

- Const Iterators for Forward Traversal

- can be obtained from all standard sequence containers either with:

- (1) member functions:

```
container.cbegin() -> @first_element
container.cend()   -> @one_behind_last_element
```

- (2) free-standing functions (C++11):

```
std::cbegin(container) -> @first_element
std::cend(container)   -> @one_behind_last_element
```

- Emptiness Query

- either with:

- (1) member function

```
container.empty() -> true, if container has no elements
```

- (2) free-standing function (C++11)

```
std::empt(container) -> true, if container has no element
```

- Type Interface

```
container::value_type
container::size_type
container::iterator
container::const_iterator
...
```

- Example

```
using con_t = std::vector<double>;
con_t::size_type i = 0;           // std::size_t
auto x = con_t::value_type{0};    // double
```

array<T, size>

- + overhead-free random access

- + fast traversal; good for linear searches
- * 'size' has to be constant expression (= known at compile time)
- * does not support size-changing operations (resize, insert, ...)
- potentially slow if element type has high copy/assignment cost (reordering elements requires copying/moving them)

- Example

```
#include <array>

std::array<int,6> a {4,8,15,16,23,42};
cout << a.size();    // 6
cout << a[0];        // 4
cout << a[3];        // 16
cout << a.front();   // 4
cout << a.back();    // 42

std::array<int,3> b {7,8,9};
a = b;               // COMPILER ERROR: types don't match
```

vector<T>

- C++'s "default" container

- + overhead-free random access
- + fast traversal; good for linear searches
- + insertion at the end in amortized constant time
- potentially slow if insert/erase operations at front and/or random positions dominate
- potentially slow if element type has high copy/assignment cost (reordering elements require copying/moving them)
- all operations that can change capacity (insert, push_back, ...) may invalidate references/pointers to any vector element
- potentially long allocation times for very large amount of values
-> can be mitigated:

https://hackingcpp.com/cpp/recipe/uninitialized_numeric_array.html

- Quick Recap

```
#include <vector>

std::vector<int> v {2,4,5};    // 2 3 5
v.push_back(6);               // 2 4 5 6
v.pop_back();                 // 2 4 5
v[1] = 3;                     // 2 3 5

cout << v[2];                  // prints 5
for (int x : v) cout << x << ' '; // prints 2 3 5

v.reserve(8);                  // 2 3 5 _ _ _ _ _
v.resize(5, 0);                 // 2 3 5 0 0 (0) _ _
cout << v.capacity();           // prints 8
cout << v.size();               // prints 5
```

.capacity() -> 8

```
+-----+
|.size() -> 5|
+-----+
| 2 | 3 | 5 | 0 | 0 | ? | ? | ? | X | contiguous buffer
|-----|-----|-----| on the HEAP
^         ^         ^
|.begin() | .end   |
+-----+-----+
| *       *       * w | STACK
```

+-----+

- Iterator Ranges
- Insert Elements

- Insert At The End (Fastest)

```
push_back(element)
```

- Example

```
vector<int> v {1,2,3};           // 1 2 3
v.push_back(4);                  // 1 2 3 4
```

- Insert Anywhere (Potentially Slow)
 - insert positions are specified with iterators:

```
.insert(@insert_pos, element)
.insert(@insert_pos, {elem1,elem2,...})
```

- Example

```
vector<int> v {1,2,3};           // 1 2 3

v.insert(begin(v)+1, 5);         // 1 5 2 3
v.insert(begin(v), {7,8});       // 7 8 1 5 2 3
v.insert(end(v), 9);             // 7 8 1 5 2 3 9
```

- Insert From Iterator Range

```
.insert(@insert_pos, @first, @last)
```

- Example

```
vector<int> y {4,5,6,7,8};
vector<int> x {1,2};
x.insert(begin(x),
         begin(y), begin(y)+3);    // 4 5 6 1 2
```

- Insert & Construct Elements In-Place (C++11)

```
.emplace_back(arg1, arg2, ...)
.emplace(@insert_pos, arg1, arg2, ...)
```

- makes a new element with forwarded constructor arguments
- Example

```
struct p2s {                      // custom
    p2d(int x_, int y_):          // <- constructor
        x{x_}, y{y_} {}         // necessary for
    int x, y;                     // emplace(_back)
};
vector<p2d> v { p2d{2,3} };       // 2 3

// insert copy
v.push_back( p2d{6,4} );          // 2 3 | 6 4

// construct in place with
// constructor arguments: 9,7
v.emplace_back(9,7);              // 2 3 | 6 4 | 9 7

// iterator to first pos (begin...)
v.emplace(begin(v), 5,8);         // 5 8 | 2 3 | 6 4 | 9 7
```

- Erase Elements

- Erase At The End (Fast)

```
vector<int> v {1,2,3,4,5,6};    // 1 2 3 4 5 6
v.pop_back();                  // 1 2 3 4 5 _
```

- Erase Everything (Fast)

```
vector<int> v {1,2,3,4,5,6};    // 1 2 3 4 5 6
v.clear();                      // _ _ _ _ _ _
```

- Erase Anywhere (Potentially Slow)

```
.erase(@position)
.erase(@range_begin, @range_end)
```

- Example

```
vector<int> v {1,2,3,4,5,6};    // 1 2 3 4 5 6
v.erase( begin(v)+2 );         // 1 2 4 5 6 _

vector<int> v {1,2,3,4,5,6};    // 1 2 3 4 5 6
v.erase( begin(v)+1, begin(v)+4 ) // 1 5 6 _ _ _
```

- NOTE:

- Erasing does not affect the capacity
- i.e., none of the vector's memory is freed

- Shrink The Capacity / Free Memory

- May work:

```
.shrink_to_fit()
```

- ISO standard does not demand that it actually shrinks
- standard library implementation might decide not to shrink

```
vector<int> v;
// add elements
// erase elements
v.shrink_to_fit();    // C++11
```

- Guaranteed to work:

- Procedure

- (1) make temporary copy -> copy does exactly fit the elements
- (2) exchange memory buffers by swapping/moving
- (3) temporary gets automatically destroyed

- Example

```
vector<int> v;
// add elements
// erase elements

// shrink: make new copy and replace v's content with it
v = vector<int>(v);    // C++11-20

// or:
v.swap( vector<int>(v) );    // C++98-20
```

- Attention After Insert/Erase!

- All iterators into a vector are invalidated if either:
 - (1) its capacity is changed or

(2) elements are moved by:

- | | |
|----------------|-----------|
| - insert | - = |
| - push_back | - assign |
| - emplace | - resize |
| - emplace_back | - reserve |
| - erase | |

- Note:

- Swapping two vector's contents does not invalidate iterators
 - Except for the 'end' iterator

- Example

```
vector<int> v {1,2,3,4,5,6}
auto i = begin(v) + 3;
```

- Dangerous

- use of iterator after insert/erase:

```
v.insert(i,8);
cout << *i;      // dangerous
```

```
v.erase(i);
cout << *i;      // dangerous
```

- Correct

- use new valid iterator returned by insert/erase
- the returned iterator refers to the original position

```
i = v.insert(i,8);
cout << *i;
```

```
i = v.erase(i);
cout << *i;
```

- Overview: Iterator Invalidating Operations

Operations	Invalidated Iterators
'const' (ro) operations	none
swap, std::swap	only end()
reserve, shrink_to_fit	if capacity changed: all else: none
push_back, emplace_back	if capacity changed: all else: only end()
insert, emplace	if capacity changed: all else: only at or after insertion point (incl. end())
resize	if capacity changed: all else: only end() and iterators to erased elements
pop_back	iterators to erased element and end()
erase	iterators to erased elements and all after them (incl. end())
clear, operator=, assign	all

```
deque<T>
```

- Double Ended Queue

```
d -> |1|2| <-> |3|4|5| <-> |6|
```

- + constant-time random access (extremely small overhead)
- + fast traversal; good for linear searches
- + good insertion and deletion performance at both ends
- + insertion does not invalidate references/pointers to elements
- potentially slow if insert/erase operations at random positions dominate
- potentially slow if element type has high copy/assignment cost (reordering elements requires copying/moving them)
- potentially long allocation times for very large amount of values; can be mitigated:
https://hackingcpp.com/cpp/recipe/uninitialized_numeric_array.html

- Example

```
#include <deque>

std::deque<int> d {0,0,0};           // 0 0 0
d.push_back(1);                     // 0 0 0 1
d.push_front(2);                     // 2 0 0 0 1

vector<int> v {3,4,5,6};             // v: 3 4 5
d.insert(begin(d),
         begin(v), end(v);           // d: 3 4 5 6 2 0 0 0 1

d.pop_front();                       // 4 5 6 2 0 0 0 1
d.erase(begin(d)+2, begin(d)+5);     // 4 5 0 0 1
```

```
list<T>
```

- Doubly-linked List

```
1 -> |1|<->|2|<->|3|<->|4|<->|5|<->|end|
      +-----+
      ^
```

- + restructuring operations don't require elements to be moved/copied (good for storing large objects with high copy/assignment cost)
- + constant-time splicing (of complete lists)
- random access only in linear time
- slow traversal due to bad memory locality

- Example

```
#include <list>

std::list<int> l {3};                // 3

l.push_back(2);                       // 3 <-> 2
l.push_front(4);                       // 4 <-> 3 <-> 2
l.splice(begin(l)+1,
         list<int>{8,4,7});           // 4 <-> 8 <-> 4 <-> 7 <-> 3 <-> 2
l.reverse();                           // 2 <-> 3 <-> 8 <-> 4 <-> 7 <-> 4
l.sort();                              // 2 <-> 3 <-> 4 <-> 4 <-> 7 <-> 8
l.unique();                            // 2 <-> 3 <-> 4 <-> 7 <-> 8
```

```
forward_list<T>
```

- Singly-linked List

```
f1 -> |1|->|2|-> ... ->|6|->|end|
```

- + uses less memory than 'list'
- + restructuring operations don't require elements to be moved/copied (good for storing large objects with high/assignment cost)
- + constant-time plicing (of complete lists)
- random access only in linear time
- slow traversal due to bad memory locality
- only forward traversal possible
- somewhat cumbersome interface due to forward-only links:
 - > no: size(), back(), push_back(), pop_back(), insert()
 - > instead: insert_after(), splice_after(), before_begin()

- Example

```
#include <forward_list>

std::forward_list<int> l {23,42,4};      // 23 -> 42 -> 4

l.insert_after(begin(l), 5);             // 23 -> 5 -> 42 -> 4
l.insert_after(before_begin(l), 88);     // 88 -> 23 -> 5 -> 42 -> 4
l.erase_after(begin(l));                 // 88 -> 5 -> 42 -> 4
```

Guidelines

- Which Sequence Container Should I Use?

- Default choice:

```
std::vector
```

- number of elements to be stored is small and known at compile time:

- yes
 - > array
- no
 - > try std::vector

- std::vector

- too slow and/or memory usage too high?
 - > measure/profile: identify bottleneck operation(s)
 - insert/erase at begin and end dominate
 - > try std::deque
 - copying elements and/or insert/erase at random position dominate
 - > try std::list
 - memory usage a little too high and reverse traversal not needed
 - > forward_list
 - still bad: look for non-standard solution with better characteristics
 - too many memory allocations and/or consumption too high
 - > .reserve(approx_expected_size) before incrementally filling vector
 - still bad: look for non-standard solution with better characteristics
 - too slow; non-linear memory accesses dominate and can't be avoided
 - look for non-standard solution with better characteristics

Sequence Views

- Views Don't Own Resources

- An object is said to be an owner of a resource (memory, file handle,

connection, ...) if it is responsible for its lifetime

std::string_view (C++17)

```
#include <string_view>
```

- Properties
 - lightweight
 - cheap on copy
 - can be passed by value
 - non-owning
 - not responsible for allocating or deleting memory
 - read-only view
 - does not allow modification of target string
 - of character range or string(-like) object
 - std::string/"literal"/...
- Primary Use Case:
 - read-only function parameters
 - avoids temporary copies
- Avoids UNnecessary Memory Allocations
 - Motivation: Read-only String Parameters
 - We don't want/expext additional copies or memory allocations for read-only parameter!
 - Traditional choice:

```
std::string const&
```

- Problematic
 - A std::string can be constructed from string literals or an iterator range to a char sequence
 - If we pass an object as function argument that is not a string itself, but something that can be used to construct a string, e.g., a string literal or an iterator range, a new temporary string object will be allocated and bound to the const reference
- string_view avoids temporary copies:

```
#include <vector>
#include <string>
#include <string_view>

void f_cref (std::string const& s) {...}
void f_view (std::string_view s) {...}

int main () {
    std::string stdStr = "Standard String";
    auto const cStr = "C-string";
    std::vector<char> v {'c','h','a','r','s','\0'};

    f_cref(stdStr);           // no copy
    f_cref(cStr);             // temp copy
    f_cref("Literal");        // temp copy
    f_cref({begin(v),end(v)}); // temp copy

    f_view(stdStr);           // no copy
    f_view(cStr);             // no copy
    f_view("Literal");        // no copy
    f_view({begin(v),end(v)}); // no copy
}
```

- String-Like Function Parameters

If You ...	Use Parameter Type
..always need a copy of the input string inside the function	std::string "pass by value"
..want read-only access - don't (always) need a copy - are using C++17/20	#include <string_view> std::string_view
..want read-only access - don't (always) need a copy - stuck with C++98/11/14	std::string const& "pass by const reference"
..want the function to modify the input string in-place (try to avoid such "output params")	std::string & "pass by (non-const) ref"

- Making string_views

- With Constructor Calls

```
std::string s = "Some Text";

// view whole string
std::string_view sv1 { s };

// view subrange
std::string_view sv2 {begin(s)+2, begin(s)+5};
std::string_view sv3 {begin(s)+2, end(s)};
```

- With Special Literal "...sv

```
using namespace std::string_view_literals;
auto literal_view = "C-String Literal"sv;
cout << literal_view;
```

- CAREFUL: View might outlive string!

```
std::string_view sv1 {std::string{"Text"}};
cout << sv1;           // string object already destroyed!

using namespace std::string_literals;
std::string_view sv2 {"std::string Literal"s};
cout << sv2;           // string object already destroyed!
```

- TIP

- Use string_view mainly as function parameter!

- string_view Interface

```
void foo (std::string_view sv) {...}
foo ("I'm sorry, Dave.");
```

sv.size()	16	(number of chars)
sv[2]	'm'	(char at index 2)
sv.front()	'I'	(first char)
sv.back()	'.'	(last char)

sv.find("r")	6 (1st match from start)	
+-----+-----+		
sv.rfind("r")	7 (1st match from end)	
+-----+-----+		
sv.find("X")	string::npos (not found - invalid index)	
+-----+-----+		
sv.substr(4,5)	string_view of "sorry"	
+-----+-----+		
sv.contains("sorry")	true (C++23)	
+-----+-----+		
sv.starts_with('I')	true (C++20)	
+-----+-----+		
sv.ends_with("Dave.")	true (C++20)	
+-----+-----+		
sv.find_first_of("ems")	2 (1st occurrence) (C++17)	
+-----+-----+		
sv.find_last_of('r')	7 (last occurrence) (C++17)	
+-----+-----+		
sv.compare("I'm sorry, Dave.")	0 (identical)	
+-----+-----+		
sv.compare("I'm sorry, Anna.")	> 0 (same length, D > A)	
+-----+-----+		
sv.compare("I'm sorry, Saul.")	< 0 (same length, D < S)	
+-----+-----+		
sv.remove_suffix(7);	-> "I'm sorry"	
+-----+-----+		
sv.remove_prefix(4);	-> "sorry"	
+-----+-----+		

std::span (C++20)

```
#include <span>
```

- Properties
 - Lightweight
 - cheap to copy
 - can be passed by value
 - Non-owning view
 - not responsible for allocating or deleting memory
 - Of a contiguous memory block
 - e.g., std::vector, std::array, ...
- Primary Use Case:
 - as a function parameter
 - container-independent access to values
- Sequences

span<int>	sequence of integers whose values can be changed	
+-----+-----+		
span<int const>	sequence of integers whose values CAN'T be changed	
+-----+-----+		
span<int,5>	sequence of exactly 5 integers (numbers of values fixed at compile time)	
+-----+-----+		

- Example

```

      w.begin()
      ^
vector<int> w { 0, 1, 2, 3, 4, 5, 6, 7, 8 };
               ^
             w.begin()+2

```

```
span<int> s {w.begin()+2, 5};           // 2 3 4 5 6
```

- As Parameter (Primary Use Case)

```
void print_ints (std::span<int const> s);
void print_chars (std::span<char const> s);
void modify_ints (std::span<int> s);
```

- Call With Container/Range

```
std::vector<int> v {1,2,3,4};
print_ints( v );
```

```
std::array<int,3> a {1,2,3};
print_ints( a );
```

```
std::string s = "SOme Text",
print_chars( s );
```

```
std::string_view sv = s;
print_chars( sv );
```

- NOTE:

- A 'span' decouples the storage strategy for sequential data from code that only needs to access the elements in the sequence, but not alter its structure

- Explicitly Making Spans

- As View of Whole Container/Range

```
std::vector<int> w {0,1,2,3,4,5,6};
std::array<int,4> a {0,1,2,3};

// auto-deduce type/length:
std::span sw1 { w };           // span<int>
std::span sa1 { a };           // span<int,4>

// explicit read-only view:
std::span sw2 { std::as_const(w) };

// with explicit type parameter:
std::span<int> sw3 { w };
std::span<int> sa2 { a };
std::span<int const> sw4 { w };

// with explicit type parameter and length:
std::span<int,4> sa3 { a };
```

- As View of Container Subsequence

```
vector<int> w {0,1,2,3,4,5,6};

std::span s1 {begin(w)+2, 4};
std::span s2 {begin(w)+2, end(w)};
```

- Size and Data Access

```
std::span<int> s = ...;

if (s.empty()) return;
if (s.size() < 1024) { ... };
```



```
// spans in range-based for loops
for (auto x : s) { ... }

// indexed access
s[0] = 8;
if (s[2] > 0) { ... }

// iterator access
auto m1 = std::min_element(s.begin(), s.end());
auto m2 = std::min_element(begin(s), end(s));
```

- Comparing Spans

```
#include <algorithm>    // std::ranges::equal

std::vector<int> v {1,2,3,4};
std::vector<int> w {1,2,3,4};

std::span sv {v};
std::span sw {w};

bool memory_same = sv.data() == sw.data();    // false
bool values_same = std::ranges::equal(sv,sw);  // true
```

- Making Spans From Spans

```
std::vector<int> v {0,1,2,3,4,5,6,7,8};
std::span s = v;

auto first3elements = s.first(3);
auto last3elements  = s.last(3);

size_t offset = 2;
size_t count = 4;

auto subs = s.subspan(offset, count);
```

Usage Guidelines

- Views As Function Parameters

- + decouple function implementations from the data representation/
container type
- + clearly communicate the intent of only reading/altering elements in
a sequence, but not modifying the underlying memory/data structure
- + make it easy to apply functions to sequence subranges
- + can almost never be dangling, because parameters outlive all function-
local variables

```
int foo (std::span<int const> in) {...}

std::vector<int> v {...};
// v will always outlive parameter 'in'!
foo(v);
foo({begin(v), 5});
```

- + a view's target cannot invalidate the memory that the view refers to
during the function's execution (unless it is done in another thread)
- + views can speed up accesses by avoiding a level of indirection:

```
vector<int> const&  --> vector<int>    --> |1|1|3|5|7|4|1|0|
reference to vector    vector object      ^   dynamic mem block
                                   |
span<int const> -----+
span object
```

- CAREFUL When Returning Views

- not always clear what object/memory the view refers to
- returned views can be (inadvertently) invalidated

- Example

```
// which parameter is the span's target?
std::span<int const>
foo (std::vector<int> const& x, std::vector<int> const& y);
```

- Example

```
// we can assume that the returned span
// refers to elements of the vector
std::span<int const>
random_subrange (std::vector<int> const& v);

// however, this is still problematic:
auto s = random_subrange(std::vector<int>{1,2,3,4});
// 's' is dangling - vector object already destroyed!
```

- Example

```
class Payments { ...
public:
    std::span<Money const> of (Customer const&) const;
    ...
};

Customer const& john = ...;
Payments pms = read_payments(file1);
auto m = pms.of(john);
pms = read_payments(file2);
// depending on the implementation of Payments
// m's target memory might no longer be valid
// after the assignment
```

- AVOID Local View Variables

- easy to produce dangling views, because we have to manually track lifetimes to ensure that no view outlives its target
- even if the memory owner is still alive, it might invalidate the memory that a view is referring to

- Example

```
std::string str1 = "Text";
std::string_view sv {str1};
if (...) {
    std::string str2 = "Text";
    sv = str2;
}
cout << sv;      // str2 already destroyed
```

- Example

```
std::string_view sv1 {"C-String Literal"};
cout << sv1;      // ok

std::string_view sv2 {std::string{"Text"}};
cout << sv2;      // string object already destroyed!

using namespace std::string_literals;
std::string_view sv3 {"std::string Literal"s};
cout << sv3;      // string object already destroyed!
```

- Memory Invalidation By Owner
 - Containers like 'vector' might allocate new memory thus invalidating all views of it:

```
std::vector<int> w {1,2,3,4,5};
std::span s {w};
w.push_back({6,7,8,9});
cout << s[0];    // w might hold new memory
```

Standard Associative Containers

Quick Overview

- Sets

(1) Ordered Sets

```
#include <set>
```

(2) Hash Sets

```
#include <unordered_set>
```

```
set<Key> / unordered_set<Key>
```

- unique orderable / hashable keys

```
std::set<int> s {9,2,8};           // {2,8,9}

s.insert(7);                       // {2,7,8,9}
s.erase(8);                       // {2,7,9}

if (s.find(7) != end(s)) {---}     // true
if (s.contains(7)) {...}          // true (C++20)

// find returns an iterator:       // {2,7,9}
auto i s.find(7);                  // ^i
if (i != end(s)) i = s.erase(i);   // {2,9}
                                   // ^i (after)
```

```
multiset<Key> / unordered_multiset<Key>
```

- multiple equivalent keys possible

```
std::multiset<int> s;              // {}
s.insert(8);                       // {8}
s.insert(7);                       // {7,8}
s.insert(2);                       // {2,7,8}
s.insert(7);                       // {2,7,7,8}
s.erase(7);                       // {2,8}
```

- Key->Value Maps

- Ordered Key->Value Maps

```
#include <map>
```

- Hashed Key->Value Maps

```
#include <unordered_map>
```

- Maps store

```
std::pair<Key const, Value>
```

- the standard library associative containers are based on nodes that are linked by pointers
 - each node stores a pair of a key and a value

```
- std::pair<First,Second>
```

- contains two values of different or same type
- Example

```
#include <utility>

std::pair<int,double> p {4, 8.15};
cout << p.first << '\n'           // 4
     << p.second << '\n';         // 8.15

// C++17 features:
std::pair p2 {1, 2.3};             // std::pair<int,double>
auto [fst,snd] = p2;               // "structured binding"
cout << fst << " " << snd << '\n'; // 1 2.3
```

```
- map<Key,Value> / unordered_map<Key,Value>
```

- unique orderable / hashable keys
- Example

```
std::map<int,std::string> m;        // {}

m.insert({2, "B"});                 // {2:B}
m.emplace(1, "A");                  // {1:A,2:B}

m[2] = "Y";                         // modify: {1:A,2:Y}
m[3] = "C";                         // insert: {1:A,2:Y,3:C}

auto i = m.find(2);                 // -> iterator
if (i != end(m))                   // if found
    cout << i->first                // 2 (key)
         << i->second;              // Y (value)

if (m.contains(2)) {...}           // true (2 found) (C++20)

m.erase(2);                         // {1:A,3:C}
auto j = m.find(3);                 //      ^j
if (j != end(m)) j = m.erase(j);   // {1:A  }
                                   //      ^j (after)

// C++17 features:
m.insert_or_assign(4, "D");         // {1:A,4:D}
m.insert_or_assign(1, "X");         // {1:X,4:D}
m.try_emplace(4, "Z");              // {1:X,4:D}
m.try_emplace(5, "E");              // {1:X,4:D,5:E}
```

```
- multimap<Key,Value> / unordered_multimap<Key,Value>
```

- multiple equivalent keys possible
- Example

```
std::multimap<int,std::string> m;   // {}

m.emplace(1, "A");                  // {1:A}
m.insert({2, "B"});                 // {1:A,2:B}
m.emplace(1, "C");                  // {1:A,1:C,2:B}
m.erase(1);                        // { 2:B}
```

- Standard Sets And Maps Are Node-Based

- Keys or key-value pairs are stored in nodes that are "linked" by pointers
- Ordered Sets/Maps
 - usually implemented as balanced binary tree
- Unordered Sets/Maps
 - implemented as hash table

Interface: How To

- Make New Sets/Maps

- Make An Empty Set/Map

```
SetOrMapType variable;
SetOrMapType variable {};

std::set<int> s1;
std::set<int> s2 {};
std::map<int,float> m1;
std::map<int,float> m2 {};
```

- CARE!

```
std::set<int> s3 ();    // THIS IS A FUNCTION DECLARATION!
```

- Make Sets From Key Lists

```
set<KeyType>{key1,key2,...}
set{key1,key2,...}      // C++17 - key type deduced

std::set<int> s1 {12};
std::set<int> s2 {3,2,1,4,5};

std::set s3 {1,2,3,4};      // set<int>      C++17
std::set s4 {1.f,2.3f,4.5f}; // set<float>   C++17
std::set s5 {1.,2.3,4.6};   // set<double>  C++17
```

- Make Sets From Key Ranges

```
set<Key>(@keys_begin,@keys_end)
set(@keys_begin,@keys_end)    // C++17 - key type deduced

std::vector<int> v {2,3,1,4};
std::set<int> s (begin(v), begin(v)+3);
```

- Inset Keys Into Sets

- Insert Single Keys

```
.insert(key) -> pair<@pos,insert_success>

std::set<int> s;          // { }
s.insert(3);              // {3}
auto r1 = s.insert(7);    // {3,7}
cout << r1.second;        // true (inserted)
cout << *r1.first;        // 7

auto r2 = s.insert(7);
cout << r2.second;        // false (NOT inserted)
                           // why? because of uniqueness
```

- Emplace: Insert And Construct Keys In-Place C++11

```
.emplace(keyArg1,keyArg2,...) -> <@pos,insert_success>
```

```
// Construct key(s) directly inside the set using key constructor
// arguments (can prevent copying of large key objects)
```

```
using VS = std::vector<std::string>;
std::set<VS> s; // { }
s.insert( VS{"a","c"} ); // {"a","c"}
s.emplace("v","w","x"); // {"a","c"}, {"v","w","x"}}
```

- Insert Ranges of Keys

```
.insert({key1,key2,...})
.insert(@keys_begin,@keys_end)
```

- Insert/Emplace With Position Hint

```
.insert(@hint,key) -> @insert_pos
.emplace_hint(@hint,keyArg1,keyArg2,...) -> @insert_pos

// Potentially faster than regular insert: amortized constant cost,
// if @hint is before/after the key's final position in the set
```

- Insert Keys+Values Into Maps

- Insert Single Key-Value Pairs

```
.insert({key,value}) -> pair<@pos,insert_success>

// Copies/moves key-value pairs into the map
// Use emplace if your key and/or value types are expensive to copy

std::map<int,std::string> m; // { }

auto r1 = m.insert({1,"a"}); // {1:"a"}
cout << r1.second; // true (inserted)
cout << *r1.first; // "a" (key @ position)

auto r2 = m.insert({1,"b"}); // {1:"a"}
cout << r2.second; // false (NOT inserted)
cout << *r2.first; // "a" (key @ position)
```

- Emplace: Insert & Construct Key-Value Pairs In-Place

```
.emplace(key,value) -> pair<@pos,insert_success> C++11
.try_emplace(key,valArg1,valArg2,...)
    -> pair<@pos,insert_success> C++17

// Constructs key-value pairs directly inside the map using
// constructor arguments (prevents copying of large value or
// key objects)

struct P2 { // custom value type
    P2 (int x_, int y_): x{x_}, y{y_} {}
    int x, y;
};

std::map<char,P2> m; // { }
m.insert( {'a', P2{1,2}} ); // {'a':{1,2}}
m.emplace('c', P2{3,4}); // {'a':{1,2}, 'c':{3,4}}
m.try_emplace('b', 5,6); // {'a':{1,2}, 'b':{5,6}, 'c':{3,4}}
m.try_emplace('a', 8,9); // {'a':{1,2}, 'b':{5,6}, 'c':{3,4}}
```

- Insert / Access / Assign

```
[key] = value (inserts key if not found)
.at(key) = value (throws std::out_of_range if key not found)
.insert_or_assign(key,value) -> pair<@pos,insert_success> C++17
```

- Insert Ranges of Key-Value Pairs

```
.insert (@kv_pairs_begin,@kv_pairs_end)
```

- Insert/Emplace With Position Hint

```
.insert(@hint,{key,value}) -> @pos
.emplace_hint(@hint,key,value) -> pair<@pos,insert_success> C++11

// Potentially faster than regular insert: amortized constant cost,
// if @hint is before/after the key's final position in the map
```

- Find/Access/Count Keys

- Check If Set/Map Contains Key

```
.find(key) != @end -> true, if container has 'key'
.contains(key) -> true, if container has 'key' C++20
```

```
std::set<int> s {1,3,5,6};

if (s.find(6) != end(s)) {...} // true
if (s.find(7) != end(s)) {...} // false

if (s.contains(6)) {...} // true
if (s.contains(7)) {...} // false
```

- Find/Get Key Position

```
.find(key) -> @position, if 'key' found or @end otherwise

std::set<int> s {1,3,5,6}; // {1, 3, 5, 6 }
auto i = s.find(3); // ^i
if (i != end(s)) { // true
    cout << *i; // 3
    i = s.erase(i); // {1, 5, 6 }
} // ^i (after)

auto j = s.find(4);
if (j != end(s)) {...} // false
```

- Access / Assign (maps only)

```
[key] = value (inserts key if not found)
.at(key) = value (throws std::out_of_range if key not found)

std::map<int,std::string> m; // { }
m[1] = "a"; // {1:"a"}
m[2] = "b"; // {1:"a", 2:"b"}
m[1] = "x"; // {1:"x", 2:"b"}

try {
    m.at(2) = "y"; // {1:"x",2:"y"}
    m.at(3) = "z"; // <- throws exception
} catch(std::out_of_range&) {...}
```

- Count Key Occurrences

```
.count(key) -> number of occurrences of 'key'
```

- Iterate Over Elements

- It is not possible to modify keys through iterators or in range based loops
 - This could break the container invariant: key ordering or position

in the hash table

- Ranged-Based 'for' Loops

- Sets:

```
for (auto lightKey : mySet) {...}
for (auto const& heavyKey : mySet) {...}
```

- Maps:

```
for (auto const& keyValuePair : myMap) {...}
for (auto const& [key,value] : myMap) {...}
```

- Examples

```
// keys that are cheap to copy
std::set<int> s1 {...};
for (auto key : s1) {...};

// keys that are expensive to copy
std::set<std::string> s2 {...};
for (auto const& key : s2) {...}

// C++17: lightweight key-value pairs
std::map<int,std::string> m2 {...};
for (auto [key,value] : m2) {
    cout << key << ":" << value << " ";
}

// C++17: expensive-to-copy key-value pairs
std::map<int,std::string> m2 {...};
for (auto const& [key,value] : m2) {
    cout << key << ":" << value << " ";
}

// C++11: lightweight key-value pairs
std::map<int,std::string> m1 {...};
for (auto kv : m1) {
    cout << kv.first << ":" << kv.second << " ";
}

// C++11: expensive-to-copy key-value pairs
std::map<int,std::string> m1 {...};
for (auto const& kv : m1) {
    cout << kv.first << ":" << kv.second << " ";
}
```

- Obtain Iterators (Forward Direction)

- Either with:

- Member functions:

```
container.begin() -> @first_element
container.end()   -> @one_behind_last_element
```

- Free-standing function:

```
std::begin(container) -> @first_element (C++11)
std::end(container)   -> @one_behind_last_element (C++11)
```

- Obtain Iterators (Reverse Direction)

- Only available for ordered containers

- Either with:

- Member functions:


```

container.rbegin() -> @last_element
container.rend()   -> @one_behind_first_element

```

- Free-standing function:

```

std::rbegin(container) -> @last_element (C++11)
std::rend(container)   -> @one_behind_first_element (C++11)

```

- Get Range of Equivalent Keys (available for all sets/maps)

```
.equal_range(key) -> pair<@range_begin,@range_end>
```

- Follows the usual iterator range convention:

- @range_end points to one behind the last element in the range

```

std::multiset<int> s {2,4,4,4,6};

// 4 is a non-unique key:           // {2,4,4,4,6}
auto e4 = s.equal_range(4);         //   ^   ^
cout << *(e4.first)                 // 4 (first in range)
    << *(e4.second);                // 6 (1 behind last)
auto n = distance(e4.first,         // n: 3 (range size)
                  e4.second);

// 1 is smaller than all:           // v
auto e1 = s.equal_range(1);         // {2,4,4,4,6}
// -> empty range                   //   ^

// 2 is the smallest key:           // v
auto e2 = s.equal_range(2);         // {2,4,4,4,6}
// -> range with 1 element          //   ^

// 3 is in between keys:           // v
auto e3 = s.equal_range(3);         // {2,4,4,4,6}
// -> empty range                   //   ^

// 6 is the largest key:           // v
auto e6 = s.equal_range(6);         // {2,4,4,4,6 }
// -> range with 1 element          //   ^

// 7 larger than all:              // v
auto e7 = s.equal_range(7);         // {2,4,4,4,6 }
// -> empty range                   //   ^

```

- Get Upper/Lower Key Bound Positions (ordered sets/maps only)

```
.lower_bound(key) -> @first_not_less_than_key
.upper_bound(key) -> @first_greater_than_key
```

```

std::multiset<int> s {2,4,4,4,6};

// 1 is smaller than all:           // {2,4,4,4,6}
auto l1 = s.lower_bound(1);         //   ^
auto u1 = s.upper_bound(1);         //   ^
if (l1 != end(s)) cout << *l1;      // true -> 2
if (u1 != end(s)) cout << *u1;      // true -> 2

// 2 is the smallest key:           // {2,4,4,4,6}
auto l2 = s.lower_bound(2);         //   ^
auto u2 = s.upper_bound(2);         //   ^
if (l2 != end(s)) cout << *l2;      // true -> 2
if (u2 != end(s)) cout << *u2;      // true -> 4

// 3 is in between keys:           // {2,4,4,4,6}
auto l3 = s.lower_bound(3);         //   ^

```

```

auto u3 = s.upper_bound(3);           //      ^
if (l3 != end(s)) cout << *l3;       // true -> 4
if (u3 != end(s)) cout << *u3;       // true -> 4

// 4 is a non-unique key:             // {2,4,4,4,6}
auto l4 = s.lower_bound(4);           //      ^
auto u4 = s.upper_bound(4);           //          ^
if (l4 != end(s)) cout << *l4;       // true -> 4
if (u4 != end(s)) cout << *u4;       // true -> 6

// 6 is is the largest key:          // {2,4,4,4,6, }
auto l6 = s.lower_bound(6);           //          ^
auto u6 = s.upper_bound(6);           //          ^
if (l6 != end(s)) cout << *l6;       // true  -> 6
if (u6 != end(s)) cout << *u6;       // false (@end)

// 7 is larger than all:              // {2,4,4,4,7, }
auto l7 = s.lower_bound(7);           //          ^
auto u7 = s.upper_bound(7);           //          ^
if (l7 != end(s)) cout << *l7;       // false (@end)
if (u7 != end(s)) cout << *u7;       // false (@end)

```

- Get Size / Query Emptiness

- Either with

- Member functions:

```

container.empty() -> true, if container has no keys
container.size()  -> total number of keys or key-value pairs

```

- Free-standing functions:

```

std::empty(container) -> true|false (C++17)
std::size(container)  -> #(keys/key-value pairs) (C++17)

```

- Example

```

std::map<int,std::string> m;           // { }
cout << m.empty();                     // true
cout << m.size();                       // 0

//m.insert(3,"x");                     // {3:"x"}
m.insert( {3,"x"} );
cout << m.empty();                     // false
cout << m.size();                       // 1

//m.insert(1,"y");                     // {1:"y",3:"x"}
m.insert( {1,"y"} );
cout << m.size();                       // 2

```

- Erase Elements/Ranges

- Erase A Single Key / Key-Value Pair

```

.erase(key) -> number of deleted keys/key-value pairs

```

- Example

```

std::multiset<int> s {1,3,3,7,9};     // {1,3,3,7,9}
auto n = s.erase(3);                  // {1,7,9}
cout << n;                             // 2

s.erase(7)                             // {1,9}

```

- Erase With Iterator (Ranges)

```
.erase(@position)          -> @behind_deleted
.erase(@range_begin,@range_end) -> @behind_last_deleted
```

- Example

```
std::set<int> s {1,3,6,8,9};          // {1,3,6,8,9}
auto i = s.find(3);                  // ^i
i = s.erase(i);                      // {1,6,8,9}
                                     // ^i (after)

auto j = s.erase(begin(s),s.find(8)); // {1,6,8,9}
                                     // ^ ^
                                     // {8,9}
                                     // ^j
```

- Erase Everything

```
std::set<int> s {1,2,5,8};          // {1,2,5,8}
s.clear();                          // { }
s.empty();                          // true
```

- Extract And (Re-)Insert Nodes (C++17)

```
.extract(key) -> node
.extract(@position) -> node
.insert(node) -> @insert_position
```

container::node_type

- Important Member Functions:

```
.key()      -> key_type&
.mapped()   -> mapped_type&
```

- Transfer Key(-Value Pair) Between Sets/Maps (without copying or moving key/value objects)

```
set<string> s {"a","b","e"};          // s: {"a","b","e"}
set<string> t {"x","z"};              // t: {"x","z"}

t.insert(s.extract("a"));             // s: {"b","e"}
                                     // t: {"a","x","z"}
```

- Change Keys Without Copying It

```
set<string> s {"a","c","e"};          // {"a","c","e"}
auto na = s.extract("a");             // get node
na.key() = "z";                       // change key
// move node back in
s.insert(std::move(na));              // {"c","e","z"}

map<int,string> m {{2,"a"},{3,"x"}};   // {2:"a",3:"x"}
auto n2 = m.extract(2);               // get node
//n2.key() = 5;                       // DOES NOT WORK
n2.value() = "z";                     // change value
// move node back in
m.insert(std::move(n2));              // {3:"x",5:"a"}
```

- Copy And Assign Entire Sets/Maps

- Copying

- Creates a new container object and copies all contained keys and values

- Assigning

- All contained objects are copied from source to assignment target

- TIP

- Copying / copy-assigning sets or maps might be expensive, if they contain a large number of elements

- Example

```
std::set<int> a {4,7,1,9};      // a: {4,7,1,9}
std::set<int> b {1,2,3};      // b: {1,2,3}
bool equal1 = (a == b);      // false

// copy assignment:
a = b;                        // a: {1,2,3}
bool equal2 = (a == b);      // true

b.clear();                    // b: {}   a: {1,2,3}
bool equal3 = (a == b);      // false

// ways of making copies:
std::set<int> a2 {a};          // a2: {1,2,3}
std::set<int> a3 (a);          // a3: {1,2,3}
std::set<int> a4 = a;          // a4: {1,2,3}
auto a5 {a};                  // a5: {1,2,3}
auto a6 (a);                  // a6: {1,2,3}
auto a7 = a;                  // a7: {1,2,3}
```

- Assigning Key Ranges

```
assign(@keys_begin, @keys_end) // DOES NOT EXIST/WORK?
```

- Compare Entire Sets/Maps

- Comparisons are value-based:

- Comparing two associative containers compares keys/key-value pairs

- Equality

- Comparison of all set/map types with == and !=
- Two associative containers are equal, if their key(-value) content is equal

- Lexicographical

- Comparison of ordered set/maps with <, <=, >, >=

- Examples

```
std::set<int> s1 {1,2,7,8,9};
std::set<int> s2 {1,4,5};

bool s1_equals_s2 = (s1 == s2); // false
bool s1_unequal_s2 = (s1 != s2); // true
bool s1_smaller_s2 = (s1 < s2); // true
bool s1_greater_s2 = (s1 > s2); // false

std::map<int,std::string> m1 {{1,"z"},{2,"f"},{7,"b"}};
std::map<int,std::string> m2 {{1,"a"},{3,"x"}};

bool m1_equals_m2 = (m1 == m2); // false
bool m1_unequal_m2 = (m1 != m2); // true
bool m1_smaller_m2 = (m1 < m2); // true
bool m1_greater_m2 = (m1 > m2); // false
```

- Merge Entire Sets/Maps (C++17)

- Integrate nodes from source into target without copying keys or values

```
target.merge(source)
```

- Runtime complexity

```
source.size() * log(target.size() + source.size())
```

- Example

```
set<string> s {2,7};
set<string> t {1,5,8,9};
t.merge(s)           // s: {} t: {1,2,5,7,8,9}
```

- Inspect/Control Hash Table (unordered sets and maps only)

- NOTE

- The exact number of hash buckets as a function of the number of inserted elements is not standardized and depends on the standard library implementation

- Check And Control Hash Table Size

```
.size()           -> #elements (#keys/key-value pairs)
.reserve(#elements)  make space for at least '#elements'
                    (might trigger rehash)
.bucket_count()    -> #hash_table_buckets
.bucket_size(bucket_index) -> #elements_in_bucket
.load_factor()     -> #non_empty_buckets / #bucket_count
.rehash(new_min_bucket_count) (also takes max_load_factor into
                             account)
```

- Example

```
unordered_map<int,string> m {{6,"x"},{4,"a"},{7,"n"},{2,"z"}};
cout
    << m.size()           // 4
    << m.bucket_count()  // 5
    << m.bucket_size(1) // 1
    << m.bucket_size(2) // 0
    << m.bucket_size(3) // 2
    << m.load_factor(); // 0.6 = 3/5
```

```


                                hash collision:
                                h(2) = h(4)  ----> +---+---+---+
                                                    | 6 | x |   |
                                                    +---+---+---+
                                +-----+---+ /
                                +-----+---+
+-----+ 4 | h(6) | * ---
| m | * --> +-----+---+
+-----+ 3 | h(4) | * -----> +---+---+---+
                                | 4 | a | * |
                                +---+---+---+
                                2 |   |   |   |
                                +-----+---+
                                +-----+---+
bucket index ---- 1 | h(7) | * ---
                                +-----+---+
                                0 |   |   |   |
                                +-----+---+
                                / hash table  -> +---+---+---+
                                /               | 7 | n |   |
                                /               +---+---+---+
                                hash bucket

```

- Example

```
unordered_map<int,string> m {{6,"x"},{4,"a"},{7,"n"},{2,"z"}};
// set to at least 7 buckets:
m.rehash(7)
cout
    << m.size()           // 4
    << m.bucket_count() // 7
```

```
<< m.bucket_size(1) // 1
<< m.bucket_size(2) // 0
<< m.load_factor(); // 0.57 = 5/7
```



- Check And Control Load Factor

```
.load_factor()      -> #occupied_slots / #all_slots
.max_load_factor() -> max_allowed_load_factor
.max_load_factor(new_maximum) (might trigger rehash)
```

- Access Or Iterate Over Hash Buckets

```
.bucket(key)          -> index_of_bucket_with_key
.bucket_size(bucket_index) -> number_of_nodes_in_bucket
```

- Get (const_)local_iterator


```
.begin(bucket_index) -> @bucket_begin
.end(bucket_index)   -> @bucket_end
```

- Get Types: Key/Value/...

```
using set_t = std::set<double>;
set_t::size_type i = 0;          // std::size_t
set_t::key_type{0} k = 0;        // double
set_t::value_type{0} v = 0;      // double

using map_t = std::map<int, std::string>;
map_t::key_type{0} k = 0;         // int
map_t::mapped_type{0} m = 0;      // string
map_t::value_type{0, "a"} v {1, "a"}; //pair<const int, string>
```

- Example



```

      key_type    mapped_type
      |-----|
      value_type

```

- Make Keys Orderable

- Key Comparison is Based on Equivalence

- a and b are equal

if a == b is true (i.e. their values are the same)

- a and b are equivalent

if !(a < b) && !(b < a) is true (i.e., neither one is ordered before the other)

- C++ standard library

- Uses equivalence based on "less than" for ordering objects (according to strict weak ordering)

- Alternative 1: Supply A Custom Comparator

- Ordered containers take an additional type parameter

```

set<Key, KeyComp>
map<Key, Mapped, KeyComp>

```

- The type 'KeyComp' needs to provide a public member function

```
bool operator() (Key const&, Key const&) const
```

which returns true, if the first argument should be ordered before the second

- The default is

```
std::less<Key>
```

- Alternative 2: Make Your Type Comparable

- You should only do that

- (1) if objects of your type can be ordered in a way that is "natural" and unambiguous
- (2) if you can at least provide a strict weak ordering

- More details:

<https://hackingcpp.com/cpp/lang/comparisons.html>

- Make Keys Hashable

- Unordered containers take an additional type parameter

```

set<Key, Hasher>          ???
map<Key, Mapped, Hasher>  ???

```

- Shouldn't it be:

```

unordered_set<Key, Hasher>
unordered_map<Key, Mapped, Hasher>

```

- The type 'Hasher' needs to provide a public member function

```
std::size_t operator() (Key const&) const
```

which returns the hash value (= an unsigned index) for a given key

- The default is

```
std::hash<Key>
```

- Example

```
struct TM_hash {
    // 32bit integer hash by T.Mueller
    constexpr std::size_t
    operator () (std::uint32_t k) const noexcept {
        k = (( k >> 16) ^ k ) * 0x45d9f3b;
        k = (( k >> 16) ^ k ) * 0x45d9f3b;
        k = (( k >> 16) ^ k );
        return k;
    }
}

// make set with custom hasher
std::unordered_set<std::uint32_t,TM_hash> s;
...
// get copy of hasher from set
auto h = s.hasher();

// custom key type
class A { ... };
// hasher function class
struct A_hash {
    std::size_t operator () (A const& k) const noexcept {
        ... // suitable hash function
    }
};

// make set with custom hasher
std::unordered_set<A,A_hash> s;
...
```

- Related

- Examples of 7 Handy Functions For Associative Containers

<https://www.cppstories.com/2021/handy-map-functions/>

- A Tour Of C++: Containers and Algorithms

<https://isocpp.org/files/papers/4-Tour-Algo-draft.pdf>

Standard Algorithms Introduction

- C++'s Standard Algorithms are

- algorithmic building blocks
- operating on (iterator) ranges of elements
- implemented as free-standing function
- generic: implemented in a (mostly) container/element-agnostic way
- many are customizable with function(object)s / lambdas
- well-tested and efficient

- First Example

```
min_element(@begin,                @end) -> @minimum
              |                      |
              |                      |
```


7	9	3	5	3	2	4	1	8
---	---	---	---	---	---	---	---	---

- Returns an iterator to the smallest element and thereby bot its position and value

```
#include <vector>
#include <algorithm> //std::min_element

std::vector<int> v {7,9,3,5,3,2,4,1,8,0};

// smallest in subrange (as shown above)
auto i = min_element(begin(v)+2, begin(v)+7);
auto min = *i           // int min = 2

// smallest in entire container
auto j = min_element(begin(v), end(v));
std::cout << *j << '\n';    //prints '0'
v.erase(j);                  // erases smallest element
```

- Organization

```
#include <algorithm>

- Non-Modifying Queries
    - finding elements / existence queries
    - minimum / maximum
    - comparing ranges of elements
    - binary search of sorted ranges

- Modifying Operations
    - copying / moving elements
    - replacing / transforming elements
    - removing elements
    - union/intersection/etc. of sorted ranged

#include <numeric>

- Operations on ranges of numbers (sums, reductions, ...)

#include <ranges>

- Composable range views, range utilities

#include <iterator>

- Iterator utilities (distance, next, ...)

#include <memory>

- Operations on uninitialized memory
```

- C++20
 - improved and easier to use versions of most standard algorithms
 - range and view adapters
 - more stringent handling of algorithm input requirements (based on "Concepts")


- Input Ranges

- Iterators

- Standard algorithms use iterators to traverse/access input elements
 - allows algorithms to be implemented independent from container types
 - eliminates the need for having one algorithm implementation per container type
 - new (third) party containers can be used with existing standard algorithm implementations

- Iterator Ranges

= pair p,q of iterators



- NOTE:

- End-of-range iterator q points one BEHIND the last element in the range

- Example: empty range

`begin(v)+2, begin(v)+2`

- Range Objects As Inputs (C++20)

`ranges::min_element(...) -> @minimum`

- Example

```
#include <vector>
#include <algorithm> // std::ranges::min_element

std::vector<int> v {3,5,3,2,4,1};
auto j = std::ranges::min_element(v);
std::cout << *j << '\n'; // prints '1'
```

- Algorithms in C++20's namespace `std::ranges`

- also accept single range objects like containers or views as inputs (before C++20: only iterator pairs)
- must be called with the full namespace ("namespace-qualified" in C++ parlance) because they can't be found by argument dependent lookup (= look up a function in the namespace of its arguments)
- A range is any object r for which `std::ranges::begin(r)` and `std::ranges::end(r)` return either valid iterators or end-of-range indicating sentinels

- Customization with Callable Parameters

- Many standard algorithms can be customized by passing a "callable entity" like a function, lambda or custom function object as parameter

```
min_element(@begin, @end) compare = o < o -> @minimum
min_element(@begin, @end, compare(o,o) -> bool) -> @minimum
```

- The second version of `min_elements` takes a callable entity as 3rd argument for comparing pairs of elements unlike the first version which uses operator <

- Example: min_element with Custom Type

```
#include <vector>
#include <algorithm>

struct P {int q; char c; };
std::vector<P> v { {2,'c'}, {1,'b'}, {3,'a'} };
```

- Compare using a function

```
// compares Ps by their 'q' member
bool less_q (P const& x, P const& y) {
    return x.q < y.q;
}

auto i = min_element(begin(v), end(v), less_q);
auto q1 = i->q;      // int   q1 = 1
auto c1 = i->c;      // char  c1 = 'b'
```

- Compare using a lambda

```
// use lambda to compare Ps by 'c'
auto j = min_element(begin(v), end(v),
    [](P const& x, P const& y) {
        return x.c < y.c;
    });

auto q2 = i->q;      // int   q2 = 3
auto c2 = i->c;      // char  c2 = 'a'
```

- Lambdas

- can be thought of as "anonymous function"
- can be defined within functions (regular C++ functions can not be nested)
- are function objects whose type is auto-generated by the compiler

- Parallel Execution (C++17)

- Most standard algorithms can be executed in parallel
 - Provide an 'execution policy' object as first argument:

```
#include <execution>
...
sort(std::execution::par, begin(v), end(v));
```

- Execution Policies And Effects

Execution Policy	Effect
std::execution::seq	parallelization and vectorization are not allowed
std::execution::unseq	may vectorize, but parallelization is not allowed (C++20)
std::execution::par	may parallelize, but vectorization is not allowed
std::execution::par_unseq	may parallelize, vectorize, or migrate computation accross threads. allows to invoke input element access functions in an unordered fashion, und unsequenced with respect to each other within each thread

- Compiler Support (min. required versions)

- GNU g++9
 - Requires TBB Library (Intel Thred Building Blocks)
 - Installation

```
$ sudo apt install libtbb-dev
```

- Link executable against TBB

```
$ g++ -std=c++17 ... -o exename -ltbb
```

- Microsoft MSVC 19.14 (VS 2017 15.7)
- NVIDIA NVC++

- Iterator / Range Categories

- Category = set of supported iterator/range object operations/guarantees
 - based on common algorithm requirements (input, output, efficiency, correctness, ...)
 - determined by the input range object or the host container providing the iterator

Sentinel (C++20)	iterator-like position specifier; usually only used for denoting the end of a range	supports == !=
Input	read access to objects; advancable to next position example: iterator that reads values from a file	supports * ++ == !=
Output	write access to objects;advancable to next position example: iterator that writes values to a file	supports * ++ == !=
Forward	read/write access; forward traversal, no random access; multi-pass guarantee: iterators to the same range can be used to access the same objects multiple times example: std::forward_list iteratr	supports * ++ == !=
BiDirectional	multi-pass guarantee, traversal in both directions (but no random access example: std::list iterator	supports * ++ -- == !=
RandomAccess	random access, but not necessarily to a contiguous memory block example: std::deque iterators	supports * [] ++ -- += -= - + == != < <= > >=
Contiguous	random access to contiguous memory example: std::vector iterators	supports * [] ++ -- += -= - + == != < <= > >=

- Error Messages

of generic algorithms can be quite confusing:

```
std::list x {3,2,8,1};
std::sort(begin(x), end(x));
```

- This does not compile because
 - (1) 'sort' requires random access iterators
 - (2) 'list' provides bi-directional iterators
- Always look for the first message that contains the word 'error'
- Algorithms in Namespace std::ranges (C++20)
 - requirements are checked at the call site using "Concepts"
 - requirements are overall more consistently specified
 - allow compiler error messages to be more helpful (still room for improvement)

- Related

<https://hackingcpp.com/cpp/std/algorithms/intro.html>

- Check Youtube Videos!

Container Traversal

- TIPS

- Try to only write loops if there is no well-tested (standard) library function/algorithm for the job to do
- Prefer non-random linear forward traversal for sequence containers like std::vector
 - > best performance due to cache and prefetching friendliness
- Reverse traversal is only supported by some standard containers

Forward Traversal

- Range-Based Loop

```
for (type variable : container)

+ works for all standard sequence and associative containers
+ container agnostic -> easy to change container type
+ no out-of-bounds access bugs possible
+ no signed/unsigned index type hassle
+ best performance when using sequence containers (due to linear access pattern); cache and prefetching friendly

* early exit possible with 'break;'

- not suited for algorithms that require random access patterns
```

- Example

```
std::vector<Type> v {...};

// read-only, type cheap to copy/or copy needed:
for (Type x : v) { cout << x; }
for (auto x : v) { cout << x; }
```

```
// read-only, type expensive to copy:
for (Type const& x : v) { cout << x; }
for (auto const& x : v) { cout << x; }

// modify values:
for (Type& x : v) { cin >> x; }
for (auto& x : v) { cin >> x; }
```

- for_each / for_each_n

- + convenient if having a function(object) to be applied to each element
- + works for all standard sequence and associative containers
- + container agnostic -> easy to change container type
- + no signed/unsigned index type hassle
- + self-documenting name

- out-of-bounds access bugs possible with iterator ranges

ranges::for_each(range, f(o)) (invokes f on each input elemnt)

- C++20
- {9,1,3,8,5} -> f(9),f(1),f(3),f(8),f(5)
- + no out-of-bounds access possible

```
#include <algorithm> // std::ranges::for_each
namespace ranges = std::ranges; // alias
```

```
Container<Type> v; ...
```

```
// read-only, type cheap to copy or copy needed:
ranges::for_each(v, [](Type x){ cout << x; });
ranges::for_each(v, [](auto x){ cout << x; });
```

```
// read-only, type expensive to copy:
ranges::for_each(v, [](Type const& x){ cout << x; });
ranges::for_each(v, [](auto const& x){ cout << x; });
```

```
// modify values:
ranges::for_each(v, [](Type& x){ cin >> x; });
ranges::for_each(v, [](auto& x){ cin >> x; });
```

for_each(@begin, @end, f(o)) (invokes f on each input elemnt)

- + can be used on subranges
- out-of-bounds access bugs possible

```
#include <algorithm> //std::for_each
```

```
// read-only, type cheap to copy or copy needed:
for_each(begin(v), end(v), [](Type x){ cout << x; });
for_each(begin(v)+2, end(v)+5, [](auto x){ cout << x; });
```

```
// read-only, type expensive to copy:
for_each(begin(v), end(v), [](Type const& x){ cout << x; });
for_each(begin(v), end(v), [](auto const& x){ cout << x; });
```

```
// modify values:
for_each(begin(v), end(v), [](Type& x){ cout << x; });
for_each(begin(v), end(v), [](auto& x){ cout << x; });
```

for_each_n(@begin, n, f(o)) (C++17) (invokes f on each input elemnt)

- + can be used on subranges
- out-of-bounds access bugs possible

- Explicit Use of Iterators

- + container agnostic -> easy to change container type
- + works for all standard sequence containers
- + no signed/unsigned index type hassle
- + possible to skip multiple elements
- out-of-bounds access bugs possible
- verbose

```
std::vector<int> v {1,2,3,4,5,6};

for (auto i = begin(v); i != end(v); ++i) { cout << *i; }
for (auto i = begin(v); i != end(v); ++i) { cin >> *i; }

// read-only - using const iterators
for (auto i = cbegin(v); i != cend(v); ++i) { cout << *i }
```

- Index-Based Loop

- + possible to skip multiple elements
- prone to out-of-bounds access bugs
- easy to write subtle bugs due to signed/unsigned index type conversions
- does not work for all sequence containers -> not easy to change container type
- making sure that loop doesn't modify elements requires more discipline
- verbose

```
std::vector<int> v {1,2,3,4,5,6};

for (std::size_t i = 0; i < v.size(); ++i) { cout << v[i]; }

// explicitly read-only
for (std::size_t i = 0; i < cv.size(); ++i) {cout << cv[i];}
```

Reverse Traversal

- Reverse Range-Based Loop (C++20)

```
for (type variable : container | std::views::reverse)
```

- + works for all bidirectional containers
- + no out-of-bounds access bugs possible
- + no signed/unsigned index type hassle
- * early exit possible with 'break;'

```
#include <ranges> // C++20

std::vector<int> v {1,2,3,4,5,6};
for (int x : v | std::views::reverse) {cout << x << '\n';}

// read-only, if type cheap to copy or copy needed:
for (auto x : v | std::views::reverse) {cout << x;}

//read-only, if type expensive to copy:
for (auto const& x : v | std::views::reverse) {...}
```

```
// modify items:
for (auto& x : v | std::views::reverse) {...}
```

- Reverse for_each/for_each_n

- + convenient if having a function(object) to be applied to each element
- + works for all bidirectional containers
- + easy to change container type
- + no signed/unsigned index type hassle
- + self-documenting name
- out-of-bounds access bugs possible with iterator ranges

- ranges::for_each(range, f(o))

```
#include <algorithm>      // std::ranges::for_each
#include <ranges>          // range views

namespace ranges = std::ranges;      // alias
namespace views = std::ranges::views; // alias

Container<Type> v;

// read-only, type cheap to copy or copy needed:
ranges::for_each(views::reverse(v), [](Type x){cout << x;});
ranges::for_each(views::reverse(v), [](auto x){cout << x;});

// read-only, type expensive to copy:
ranges::for_each(views::reverse(v), [](Type const& x){
    cout << x; });
ranges::for_each(views::reverse(v), [](auto const& x){
    cout << x; });

// modify values:
ranges::for_each(views::reverse(v), [](Type& x){
    cout << x; });
ranges::for_each(views::reverse(v), [](auto& x){
    cout << x; });
```

- for_each(@begin, @end, f(o))

- + can be used on subranges
- out-of-bounds access bugs possible

```
for_each(rbegin(v), rend(v), [](Type x){...});
...
```

- for_each_n(@begin, n, f(o))

- + can be used on subranges
- out-of-bounds access bugs possible

```
for_each_n(rbegin(v), 2, [](Type x){...});
...
```

- Explicit Use of Reverse Iterators

- + works for all bidirectional containers
- + no signed/unsigned index type hassle
- + possible to skip multiple elements
- out-of-bounds access bugs possible
- verbose

```
for (auto i = rbegin(v); i != rend(v); ++i) {...}
```



```
// read-only - using const iterators
for (auto i = crbegin(v); i != crend(v); ++i) {...}
```

- Reverse Index-Based Loop

- prone to out-of-bounds access bugs
- easy to write subtle bugs due to unsigned size type: implicit conversions to signed int, overflow/wrap-around, ...
- making sure that loop doesn't modify elements requires more discipline
- verbose

```
// std containers use UNsigned size types
// -> be careful not to decrement unsigned '0'
for (auto i = v.size(); i > 0; --i) { cout << v[i-1]; }

// explicitly read-only
const auto& cv = v;
for (auto i = cv.size(); i > 0; --i) { cout << cv[i-1]; }
```

Utilities

- Get Next/Previous Iterator

```
#include <iterator>
```

- std::prev and std::next

- Functions
- Provide a universal way of incrementing/decrementing iterators
 - Even if the iterator type does not support random access (e.g., 'it += 5')
- WARNING
 - Be aware that advancing non-random access iterators (e.g., those from std::list) by N steps might be costly
 - i.e., involve on the order of N memory operations

```
next(@position)          -> @one_after
next(@position, steps)   -> @steps_after
```

```
std::vector<int> v {1,2,3,4,5,6}; // 1 2 3 4 5 6 _
auto i = next(v.begin());         //   ^i
auto j = next(i, 3);              //   ^i   ^j
```

```
prev(@position)          -> @one_before
prev(@position, steps)   -> @steps_before
```

```
std::vector<int> v {1,2,3,4,5,6}; // 1 2 3 4 5 6 _
auto i = prev(v.end());           //           ^i
i = prev(i);                      //           ^i
auto j = prev(i, 3);              //   ^j   ^i
```

Standard Library min/max Algorithms

- min

```
min(a, b) -> a if (a < b) is true, b otherwise
min(a, b, cmp(o,o)->bool) -> a if cmp(a,b) is true, b otherwise
```

- Example

```
int const a = 2;
int const b = 9;
int x = std::min(a,b); // int x = 2
```

- Example

```
struct P { int q; char c; };
P pmin = std::min(p{1,'y'}, P{2,'x'}, [](P p1, P p2){
    return p1.q < p2.q }); // P min {1,'y'}
```

```
min({v1,v2,v3,...}) -> smallest_value (C++11)
min({v1,v2,v3,...}, cmp(o,o)->bool) -> smallest_value
```

- The second version uses cmp for comparing elements, while the first version uses 'operator <'

- CARE

- ALL elements in the input list {...} must have the same type!

- Example

```
int const a = 2;
int const b = 9;
int x = std::min({3,4,b,3,a,8});
```

- Example

```
std::set<int> s1 {3,5,2};
std::set<int> s2 {9,3,1};
std::set<int> s3 {4,2,6};

int s_min = std::min({s1,s2,s3});
```

- Example

```
struct P { int q; char c; };

P px {3,'x'};
P py {2,'y'};
P pz {1,'z'};

P p_min = std::min({px,py,pz},
    [](P p1, P p2){ return p1.q < p2.q; });
```

```
ranges::min(range) -> smallest_value (C++20)
ranges::min(range, cmp(o,o)->bool) -> smallest_value
```

- Returns (a const reference to) the smallest element in range
- The second version uses cmp for comparing elements, while the first version uses 'operator <'

- Example

```
std::vector<int> v {7,9,3,5,3,1,5,8};
auto x = std::ranges::min(v); // int x = 1
```

- Example

```
struct P { int q; char c; };

std::vector<P> const w {P{3,'a'},P{1,'c'},P{2,'b'}};
auto pmin = std::ranges::min(w,
    [](P const& p1, P const& p2){ return p1.q < p2.q; });
```

- max

```
max(a, b) -> a if (a < b) is false, b otherwise
max(a, b, cmp(o,o)->bool) -> a if cmp(a,b) is false, b otherwise
```

```
max({v1,v2,v3,...}) -> largest_value (C++11)
max({v1,v2,v3,...}, cmp(o,o)->bool) -> largest_value

ranges::max(range) -> largest_value (C++20)
ranges::max(range, cmp(o,o)->bool) -> largest_value
```

- minmax

```
minmax(a, b) -> {smallest,largest} (C++11)
minmax(a, b, cmp(o.o)-> bool) -> {smallest, largest}
```

- Comparison function/object cmp(a,b) must return true if 'a' should be ordered before 'b'
- Example

```
int a = 2;
int b = 9;

auto p = std::minmax(a,b); // std::pair<int,int> p {2,9}
auto min = p.first;       // int min = 2
auto max = p.second;      // int max = 9

auto [lo,hi] = std::minmax(a,b); // int lo = 2, hi = 9 (C++17)
```

```
minmax({v1,v2,v3}) -> {smallest,largest}
minmax({v1,v2,v3}, cmp(o.o)->bool) -> {smallest,largest}
```

- The second version uses cmp for comparing elements, while the first version uses 'operator <'
- CARE
 - All elements in the input list { ... } must have the same type

```
auto p = std::minxmax({3,0,b,3,a,8}); // std::pair<int,int> p
                                         // {0,9}

auto min = p.first;
auto max = p.second;

auto [lo,hi] = std::minmax({3,0,b,3,a,8});
                                         // int lo = 0, hi = 9 (C++17)
```

```
ranges::minmax(range) -> {smallest,largest} (C++20)
ranges::minmax(range, cmp(o.o)->bool) -> {smallest,largest}
```

- Returns a pair of (const references to) the smallest and largest elements in range
- The second version uses cmp for comparing elements, while the first version uses 'operator <'

```
std::vector<int> v {7,9,3,6,3,1,4,8};
auto p = std::ranges::minmax(v); // std::pair<int,int> p {1,9}

struct P { int q; char c; };
std::vector<P> const w {P{3,'a'},P{2,'b'},P{1,'c'}};
auto [lo,hi] = std::ranges::minmax(w,
    [](P p1, P p2){ return p1.q < p2.q; });
```

- clamp (C++17)

```
clamp(value, lo, hi) -> clamped_value
clamp(value, lo, hi cmp(o.o)->bool) -> clamped_value
```

- clamps value in the interval given by lo and hi
- The second version uses cmp to compare values instead of 'operator <'

```
int a = std::clamp( 8, 1, 5); // int a = 5
```

```
int b = std::clamp(-4, 1, 5); // int b = 1
int c = std::clamp(-4, -2, 5); // int c = -2
```

- min_element

```
min_element(@begin, @end)                -> @minimum
min_element(@begin, @end, compare(o.o)->bool) -> @minimum
```

```

      @begin      @end
        |          |
        v          v
    |7|9|3|5|3|2|4|1|8|0
        ^
        |
      @minimum

```

- The second version uses comp for comparing elements, while the first version uses 'operation <'

```
std::vector<int> v {7,9,3,5,3,2,4,1,8,0};

// smallest in subrange (as shown above)
auto i = min_element(begin(v)+2, begin(v)+7);
auto min = *i; // int min = 2

// smallest in entire vector
auto j = min_element(begin(v), end(v));
std::cout << *j; // print '0'

// index of smallest
auto argmin = distance(begin(v), j); // int argmin = 9

// erase at i's position
i = v.erase(i); // 7 9 3 5 3 2 4 1 8 0
//                  ^

std::cout << *i; // 7 9 3 5 3 4 1 8 0
// prints '4'    ^

ranges::min_element(v)                -> @minimum (C++20)
ranges::min_element(v, comp(o.o)->bool) -> @minimum
```

- The second version uses comp for comparing elements, while the first version uses 'operation <'

```
std::vector<int> v {7,9,8,3,6,4,0,4};

auto i = std::ranges::min_element(v);
auto min = *i; // int min = 0;
```

- max_element

```
max_element(@begin, @end)                -> @maximum
max_element(@begin, @end, comp(o.o)->bool) -> @maximum

ranges::max_element(v)                -> @maximum (C++20)
ranges::max_element(v, comp(o.o)->bool) -> @maximum
```

- minmax_element

```
minmax_element(@begin, @end)                -> {@min, @max}
minmax_element(@begin, @end, comp(o.o)->bool) -> {@min, @max}

ranges::minmax_element(v) -> {@min, @max} (C++20)
```

```

    auto [min,max] = std::ranges::minmax_element(v);

    std::cout << "min: " << *min << '\n'
               << "max: " << *max << '\n';

    ranges::minmax_element(v, comp(o.o)->bool) -> {@min,@max}

```

Standard Library Existence Queries

- any_of / all_of / none_of (C++11)

```

all_of (@begin, @end, check(o)->bool)    -> true,
any_of (@begin, @end, check(o)->bool)    -> if check yields true for all,
none_of(@begin, @end, check(o)->bool)    -> any, or none of the elements
                                           in input range

```

```

    auto const check = [](int x) { return x >= 1; };

    cout << all_of (begin(v), end(v), check);

    ranges::all_of (range, check(o)->bool) -> bool
    ranges::any_of (range, check(o)->bool) -> bool
    ranges::none_of(range, check(o)->bool) -> bool

```

- count

```

count(@begin, @end, value) -> number of occurrences

ranges::count(range, value) -> number of occurrences (C++20)

```

- count_if

```

count_if(@begin, @end, f(o)->bool) -> #elements(f=true)

    std::vector<int> v {5,4,9,1,3,2,5,6,8,9};

    auto const is_even = [](int x) { return !(x & 1); };

    auto n = count_if (begin(v)+1, begin(v)+8, is_even);    // 3
    auto m = count_if (begin(v), end(v), is_even);

    count_if(range, f(o)->bool) -> #elements(f=true)

```

Standard Library Finding Algorithms

Find / Locate One Element

- find

```

find(@begin, @end, value)    -> @1st element equal to value
                             -> @end if no match
ranges::find(range, value) (C++20)

```

- find_if

```
find_if(@begin, @end, f(o)->bool)    -> @1st_element for which f is true
                                     -> @end if no such element found
```

```
auto const f = [](int x){return x >= 6; };
```

```
auto i = find_if(begin(v)+2, begin(v)+7, f);
```

```
ranges::find_if(range, f(o)->bool) (C++20)
```

- find_if_not

```
find_if_not(@begin, @end, f(o)->bool)    -> @1st_element for which f is
ranges::find_if_not (range, f(o)->bool) -> false - @end if no such elem
```

- find_last / _if / _if_not

```
ranges::find_last(range, value) -> last2end_view (C++23)
```

```

      range
+-----+
| 2 | 1 | 7 | 1 | 1 | 5 | 8 |
+-----+
      last2end_view (empty if nothing found)
```

```
std::vector<int> v {2,1,7,1,1,5,8};
```

```
auto const result = std::ranges::find_last(v, 1);
```

```
if (not result.empty()){                // if found
    auto const value = result.front();    // int value = 1
    auto const index = distance(begin(v), begin(result)); // 4
}
```

```
for (int x : result) { cout << x << ' '; } // 1 5 8
```

```
ranges::find_last_if (range, value) -> last2end_view (C++23)
```

```
ranges::find_last_if_not(range, value) -> last2end_view (C++23)
```

```
auto const f = [](int x){ return x >= 2; };
auto const result = ranges::find_last_if_not(v,f);
```

- find_first_of

```
find_first_of(@s_begin,@s_end,@w_begin,@w_end) -> @1st match
                                                  -> @s_end if no match
```

```

@s_begin|   search here   |@s_end
+-----+
| 0 | 1 | 3 | 2 | 5 | 7 | 4 | 8 | 9 | 9 |
+-----+
      | 1 | 4 | 6 | 5 | 8 | 7 |
@w_begin|   |@w_end
+-----+
find any of these values
```

```
std::vector<int> s {0,1,3,2,5,7,4,8,9,9};
```

```
std::vector<int> w {1,4,6,5,8,7};
```

```
auto i = find_first_of(begin(s)+1,begin(s)+9,
                       begin(w)+1,begin(w)+4);
```

```
if (i != begin(s)+9) {                // true, found one
    auto const value = *i;              // int value = 5
    auto const index = distance(begin(s), i); // index = 4
```

```

    }

    ranges::find_first_of(range_s, range_w) -> @1st match
                                              -> @end_s if no match

    std::vector<int> s {3,2,5,7,4,8};
    std::vector<int> w {4,6,5};

    auto i = std::ranges::find_first_of(s,w);

```

Find Subrange in Range

- search

```

search(@s_begin,@s_end,@w_begin,@w_end)    -> @1st occurrence of range
                                              'w' inside range 's'
                                              -> @s_end otherwise

```

```

                seach here
            +-----+
@s_begin | 0 | 4 | 6 | 5 | 1 | 4 | 6 | 5 | 8 | 9 | | @w_end
          |0|4|6|5|1|4|6|5|8|9|
          |1|4|6|5|8|9|
@w_begin | 1 | 4 | 6 | 5 | 8 | 9 | | @w_end
          +-----+
          find this range

```

```

std::vector<int> s {0,4,6,5,1,4,6,5,8,9};
std::vector<int> w {1,4,6,5,8,9};

auto i search = search(begin(s)+1,begin(s)+9,
                       begin(w)+1,begin(w)+4);

```

```

ranges::search(range_s,range_w) -> subrange_view (C++20)

```

```

std::vector<int> s {1,4,6,5,8,4,6,5};
std::vector<int> w {4,6,5};

auto r = std::ranges::search(s,w);
if (not empty(r)) {
    for (int x : r) {cout << x << ' ';}      // 4 6 5
}

```

- find_end

```

find_end(@s_begin,@s_end,@w_begin,@w_end)  -> @last occurrence of
                                              range w inside range s
                                              -> @s_end otherwise

```

```

                seach here
            +-----+
@s_begin | 0 | 4 | 6 | 5 | 1 | 4 | 6 | 5 | 8 | 9 | | @w_end
          |0|4|6|5|1|4|6|5|8|9|
          |1|4|6|5|8|9|
@w_begin | 1 | 4 | 6 | 5 | 8 | 9 | | @w_end
          +-----+
          find this range

```


- mismatch

```

mismatch(@begin1, @end1, @begin2)          -> {@mismatch in range1,
mismatch(@begin1, @end1, @begin2, @end2)    @mismatch in range2}

ranges::mismatch(range1, range2)           -> {@in1, @in2}

```

- lexicographical_compare

```

lexicographical_compare(@begin1, @end1, @begin2, @end2)
-> true, if range 1 should be ordered before range 2

```

```

ranges::lexicographical_compare(range1, range2)
-> true, if range 1 should be ordered before range 2

```

```

std::vector<char> range1 = {'a','l','g','o'};
std::vector<char> range2 = {'b','c','e'};

// true:
cout << std::ranges::lexicographical_compare(range1, range2);
// false:
cout << std::ranges::lexicographical_compare(
    range1, range2, std::greater<>{});

```

- lexicographical_compare_three_way

```

lexicographical_compare_three_way(@begin1, @end1, @begin2, @end2) (C++20)
-> 3-way comparison result

```

```

result < 0 -> range1 before    range2
result = 0 -> range1 equiv. to range2
result > 0 -> range1 after     range2

```

- Comparing entire strings with C++20's 'spaceship' operator:

```

std::string r1 = "xalgori";
std::string r2 = "abced";

auto const lcB = r1 <=> r2;
cout
    << std::boolalpha
    << (lcB < 0)          // false
    << (lcB == 0)         // false
    << (lcB > 0);         // true

```